

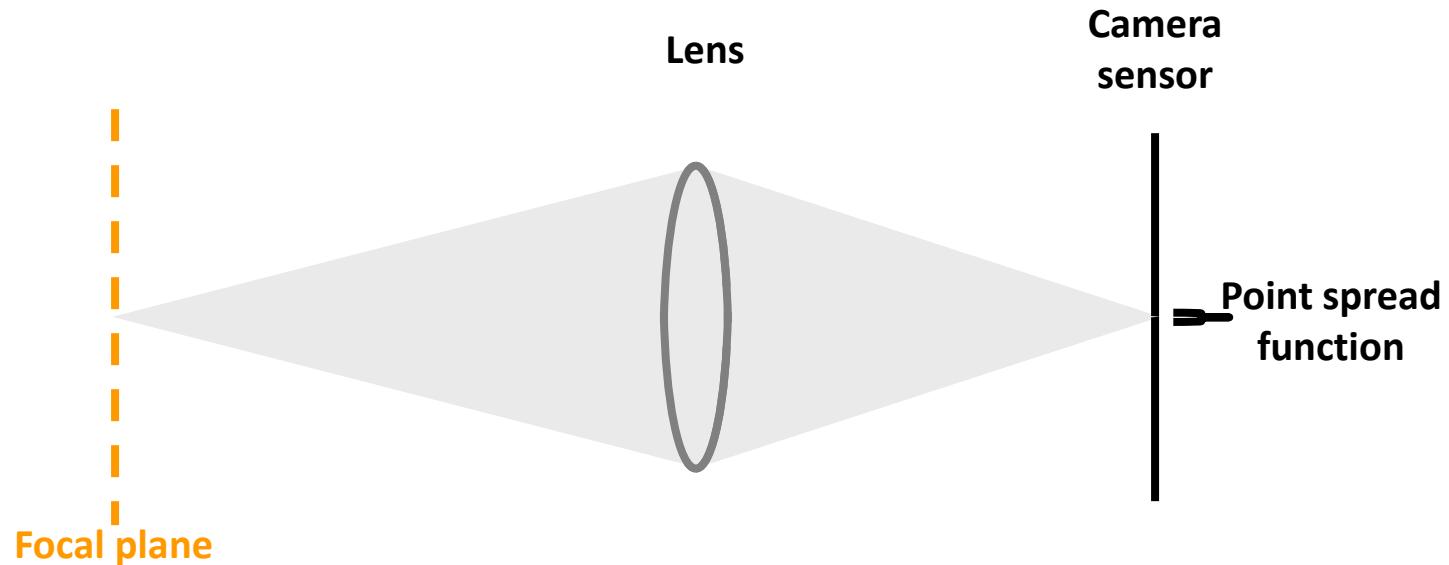
Depth from Focus PA#1

Supplementary Materials

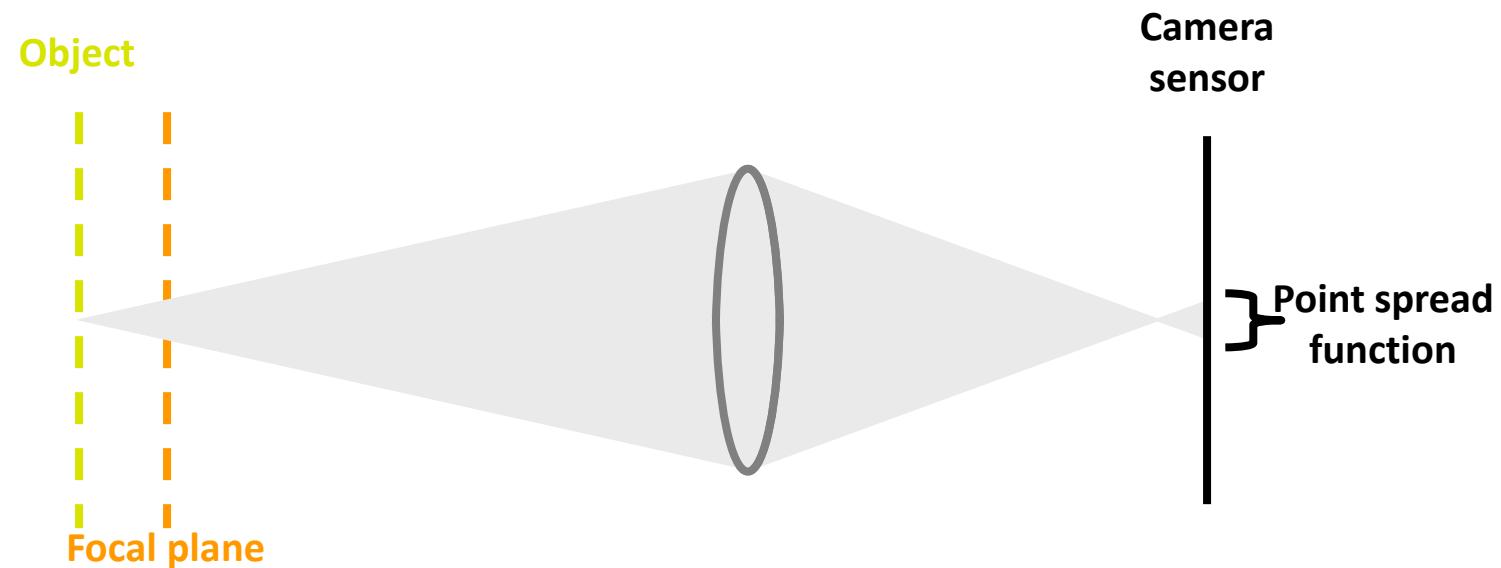
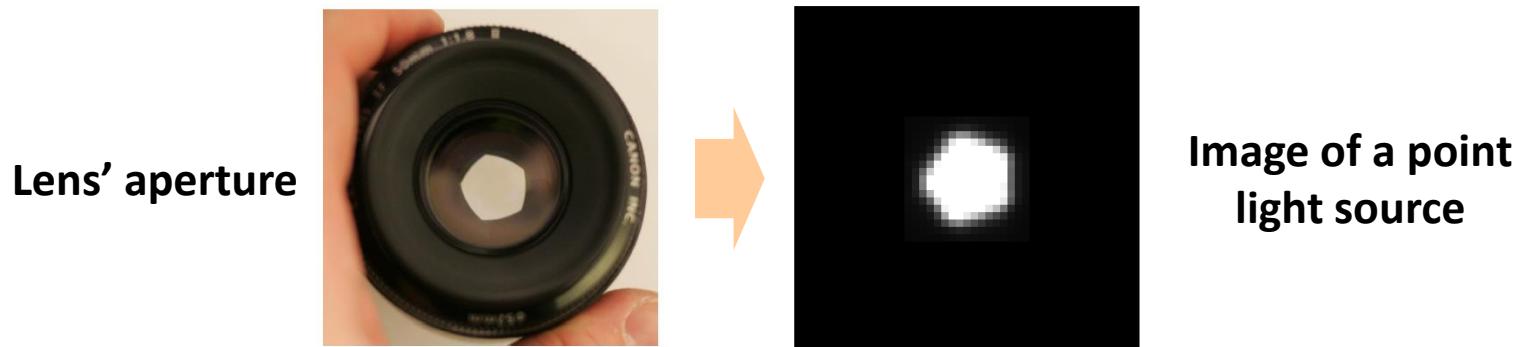
Supplementary Materials for PA#1

- Lens and Focusing
- Depth from Focus
- Graph cut optimization
- Learning based DFF method

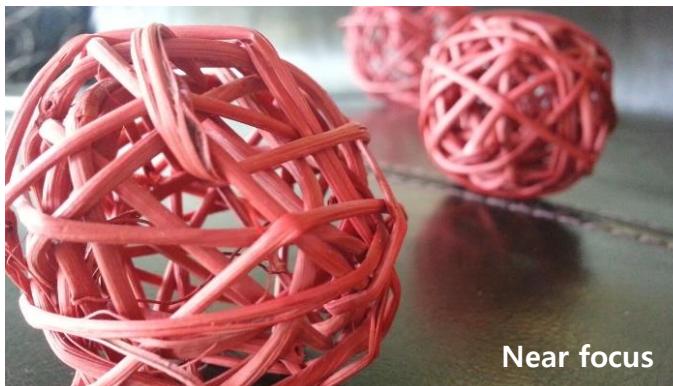
Depth and Focus



Depth and Focus



Depth and Focus



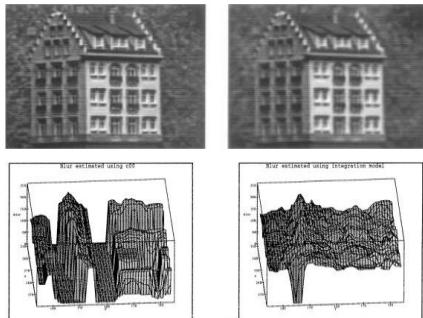
Focusing (Defocus)

- The amount of defocus is one of the priceless information that can be obtained from a single image.
- Estimated focus can be applied to many computer vision tasks.



Focusing: Applications

- Defocus information can be a cornerstone for various applications.



Depth from defocus [1]

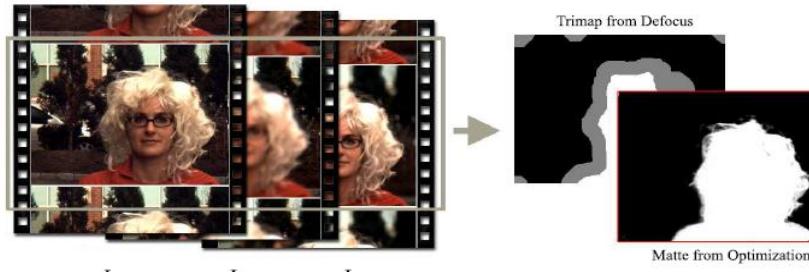
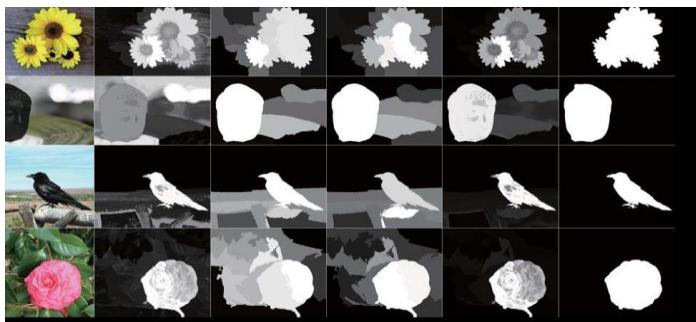


Image matting [2]



Salient region detection [3]



Digital refocusing [4]

[1] Ziou, Djemel, and Francois Deschênes. "Depth from defocus estimation in spatial domain." CVIU 2001.

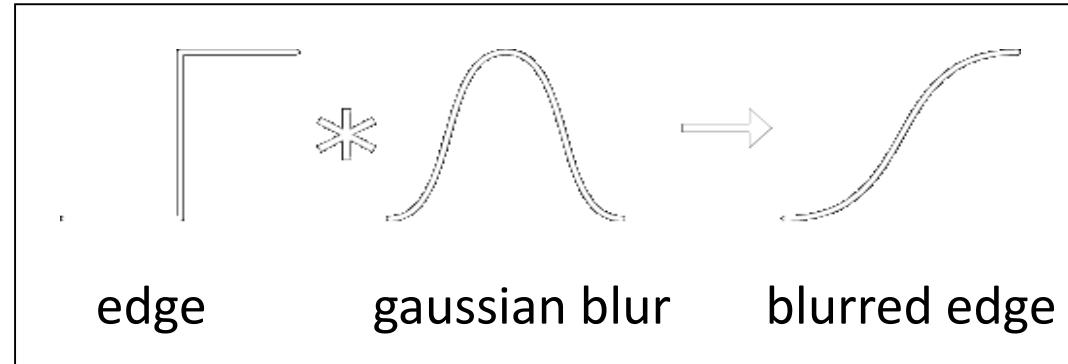
[2] McGuire, Morgan, et al. "Defocus video matting." TOG 2005.

[3] Jiang, Peng, et al. "Salient region detection by ufo: Uniqueness, focusness and objectness." ICCV 2013.

[4] Cao, Yang, Shuai Fang, and Zengfu Wang. "Digital multi-focusing from a single photograph taken with an uncalibrated conventional camera." TIP 2013

Focus Estimation at Edges

- An edge
 - a step function in intensity
- The blur of this edge
 - a Gaussian blurring kernel

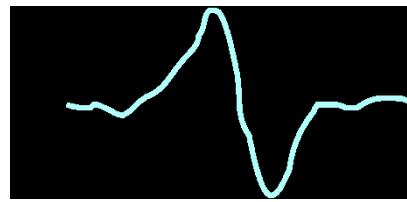


Focus Estimation at Edges

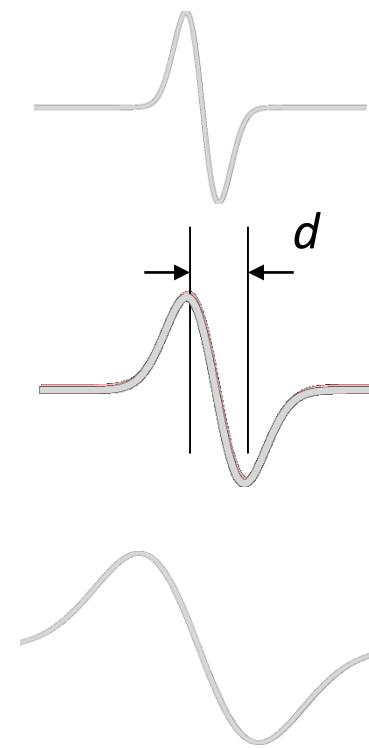
- Fit response models of various sizes



edge



2nd derivative



response model

less blurry



more blurry

Depth and Focus Measure

$$\text{Laplacian} = \frac{\partial^2 g(x, y)}{\partial^2 x} + \frac{\partial^2 g(x, y)}{\partial^2 y}$$

g is an image frame
 (x, y) is a pixel index

$$\text{Modified Laplacian (ML)} = \left(\frac{\partial^2 g(x, y)}{\partial^2 x} \right)^2 + \left(\frac{\partial^2 g(x, y)}{\partial^2 y} \right)^2$$

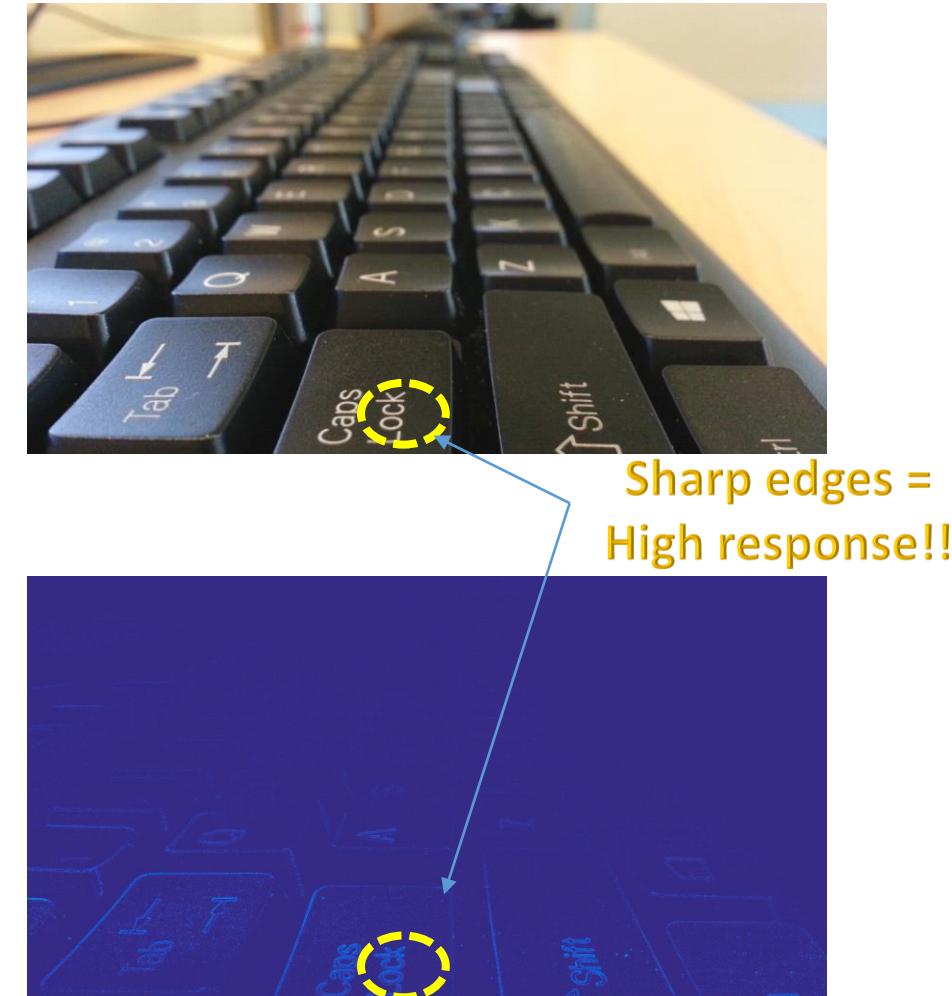
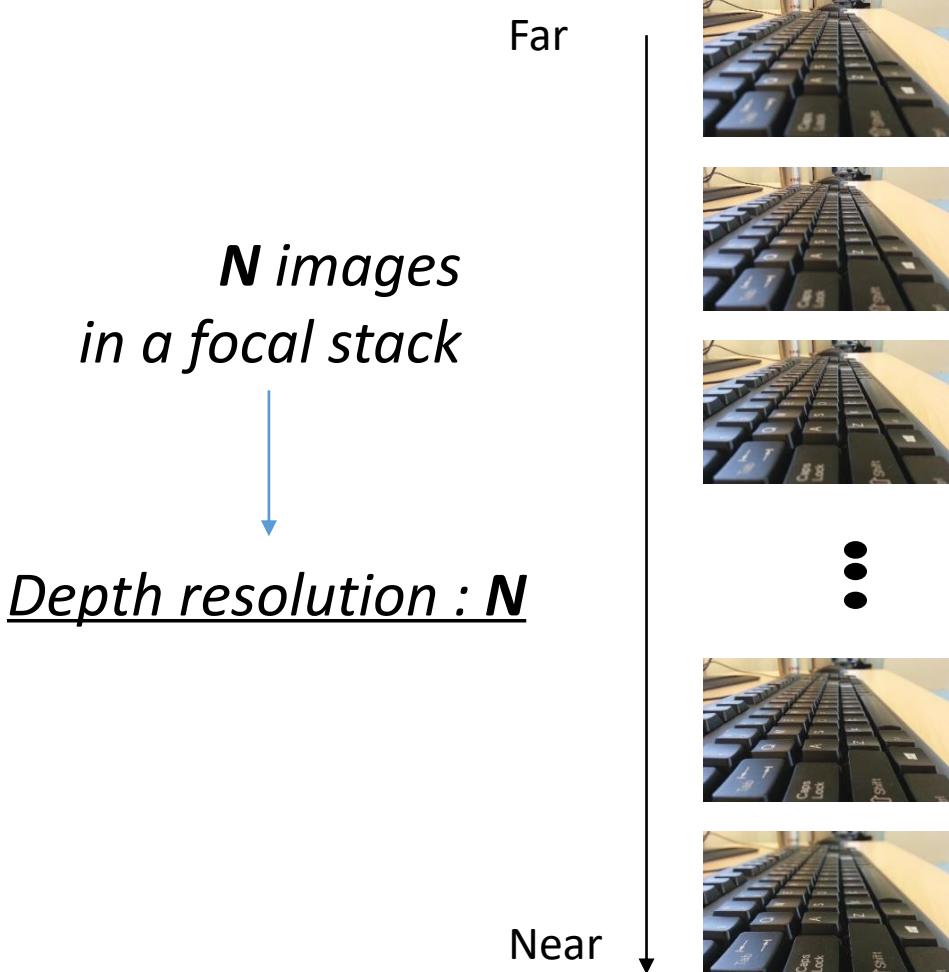
$$\text{Sum of ML} = \sum_{p(x,y)=U(x_0,y_0)} \left(\frac{\partial^2 g(x, y)}{\partial^2 x} \right)^2 + \left(\frac{\partial^2 g(x, y)}{\partial^2 y} \right)^2$$

$p(x, y)$ is a pixel in the neighborhood $U(x_0, y_0)$ of pixel (x_0, y_0)

$$\text{Tenenbaum Focus Measure} = \sum_{p(x,y)=U(x_0,y_0)} (g_x(x, y)^2 + g_y(x, y)^2)^2$$

Depth and Focus Measure

- Finding the maximum focus measure = Estimating the depth



Focal Stack to Focus Measure Volume

Far



⋮

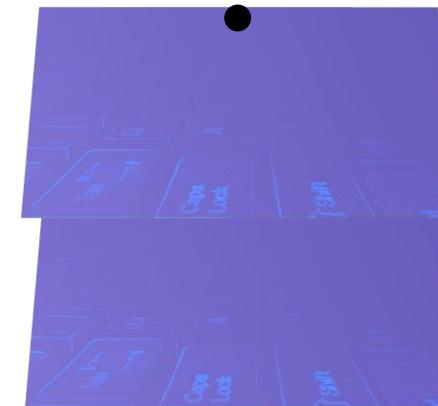
Near



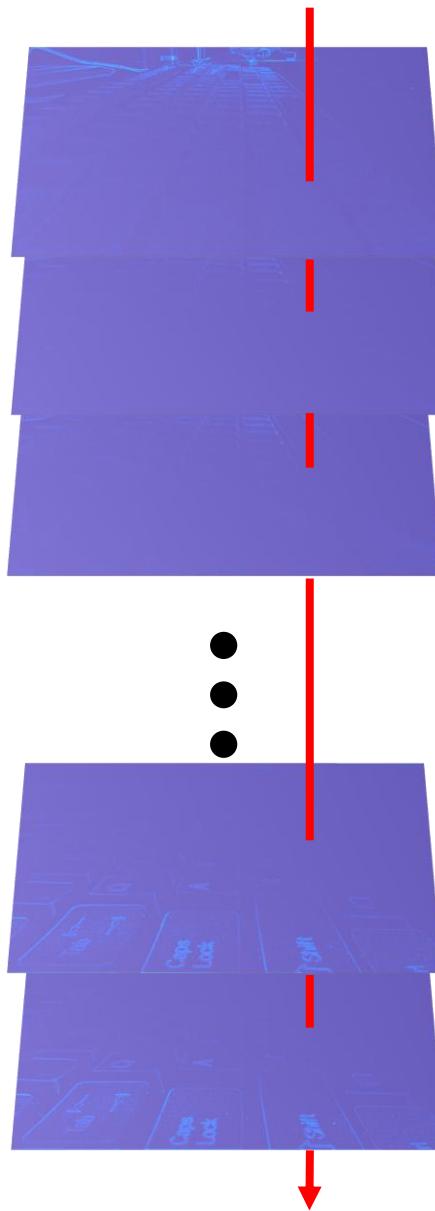
Focal stack



⋮



Focus Measure Volume to Index Map

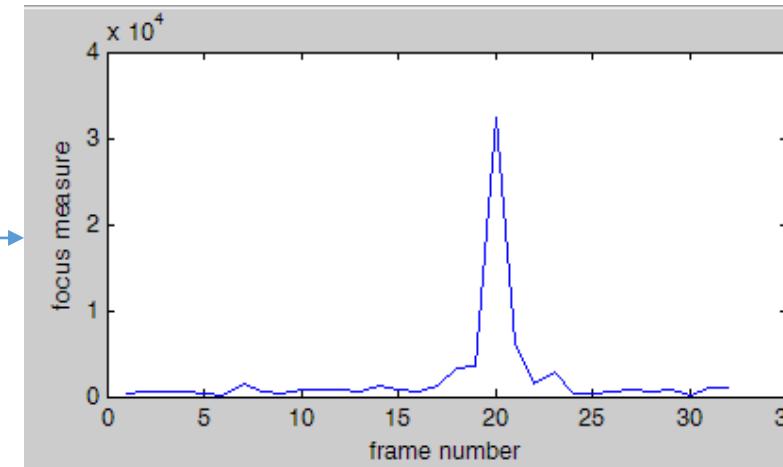


N-th

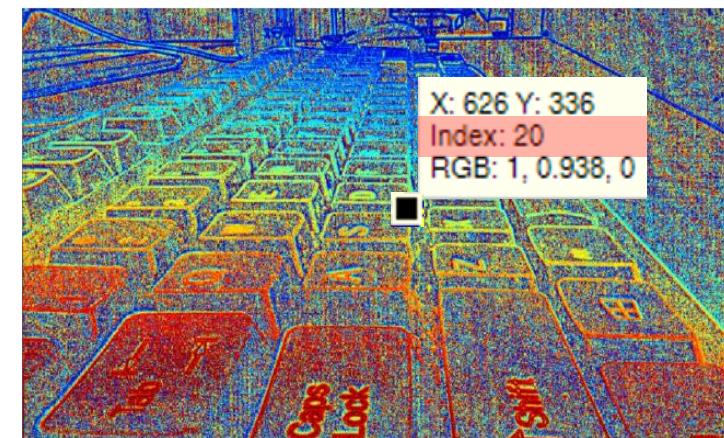
20-th

2-nd

1-st



→ Vote for the frame index with highest response.



Initial Index Map

N

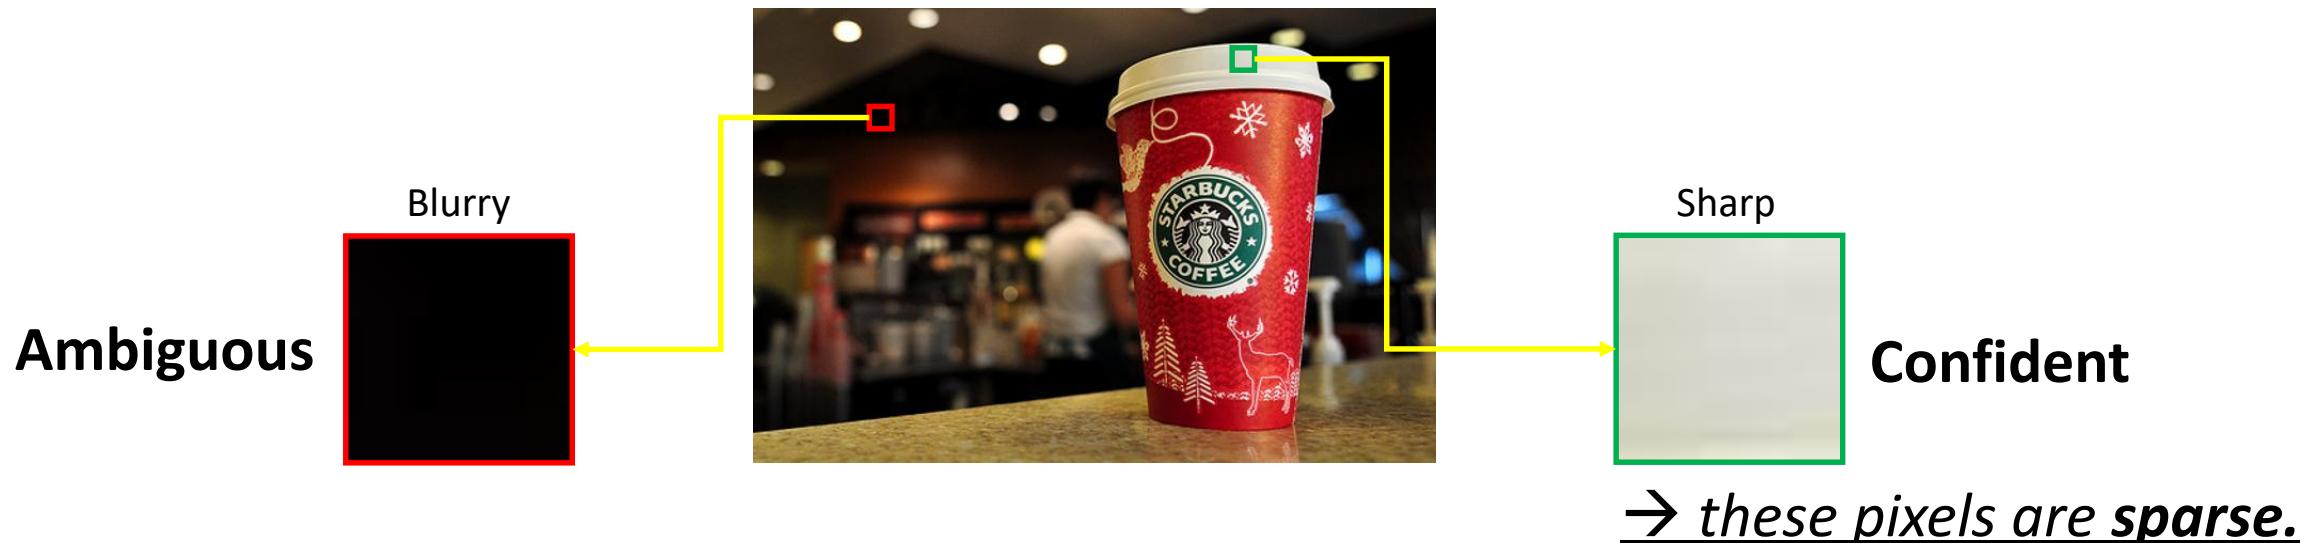
20

1

N
= # of images in a volume
= Depth resolution

Confidence of Focus Measure

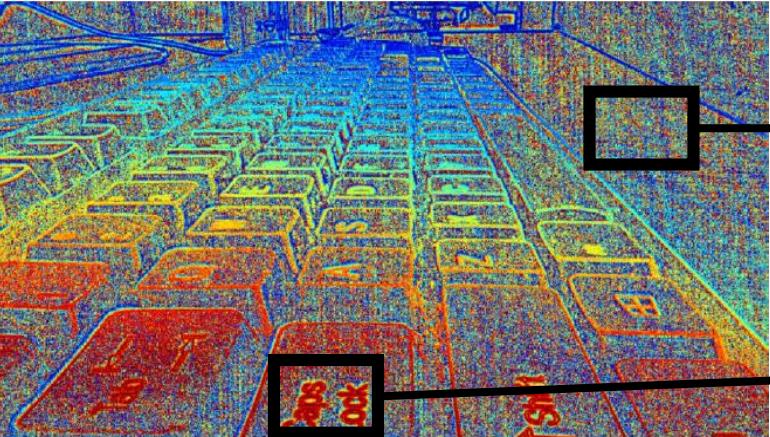
- However, focus estimation is a very challenging task.



- There exist ambiguities in homogeneous regions.
- Homogeneous regions have almost no difference in appearance when they are blurred or not.

Confidence of Focus Measure

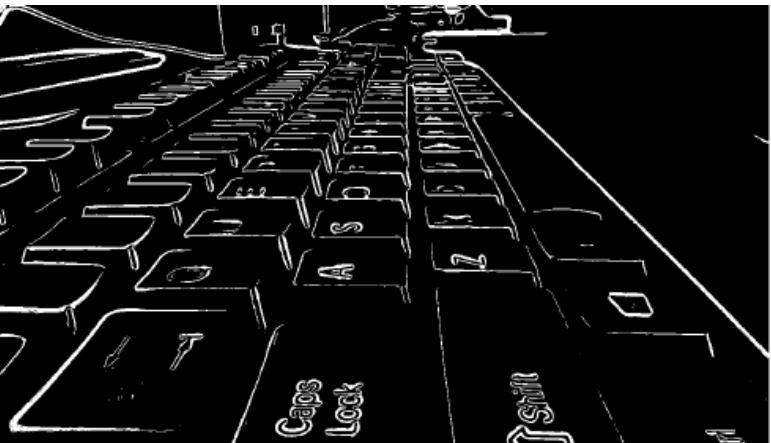
- Initial index depth map



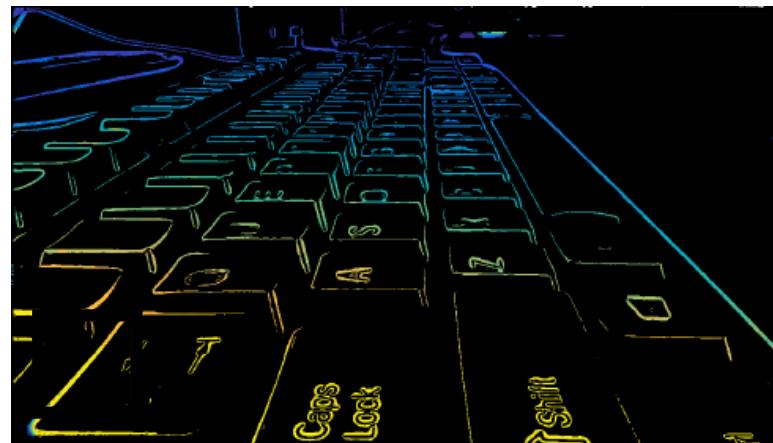
Ambiguous:
@ homogeneous regions,
Noisy, misclassified.

Confident:
@ edge regions.

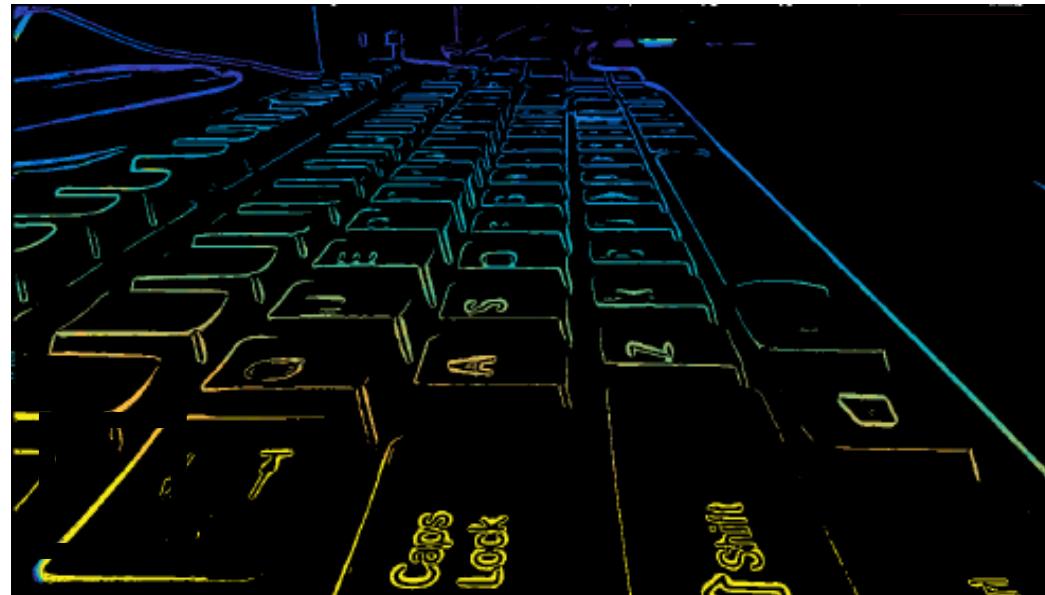
- Collect edge detections from all frames



- Sparse depth map – confident (valid) pixels

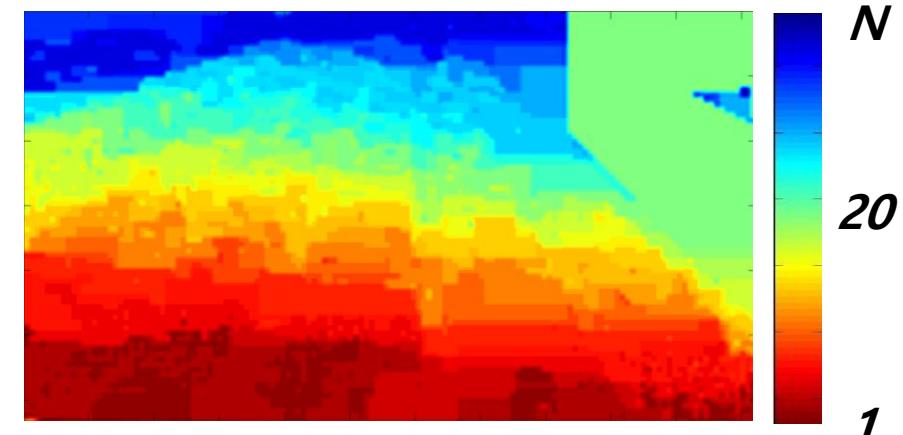


Sparse-to-Dense Depth Map



Sparse depth map – confident (valid) pixels

How to obtain dense depth map?



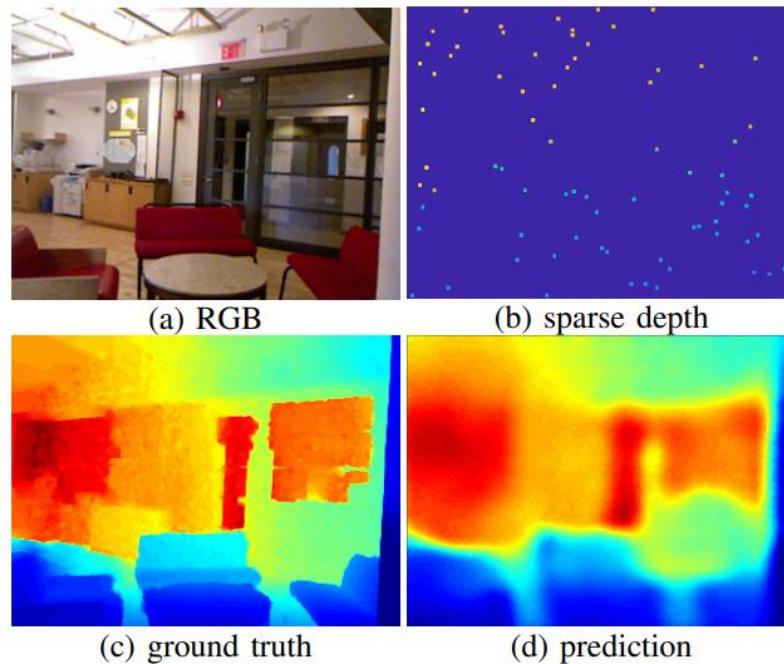
Sparse-to-Dense Propagation in Computer Vision

- Depth map propagation (from sparse sensor data points)

Sparse-to-Dense: Depth Prediction from Sparse Depth Samples and a Single Image

Fangchang Ma¹ and Sertac Karaman¹

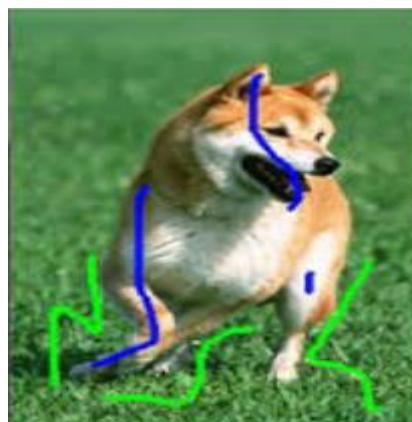
Abstract— We consider the problem of dense depth prediction from a sparse set of depth measurements and a single RGB image. Since depth estimation from monocular images alone is inherently ambiguous and unreliable, to attain a higher level of robustness and accuracy, we introduce additional sparse depth samples, which are either acquired with a low-resolution depth sensor or computed via visual Simultaneous Localization and Mapping (SLAM) algorithms. We propose the use of a single deep regression network to learn directly from the RGB-D raw data, and explore the impact of number of depth samples on prediction accuracy. Our experiments show that, compared to using only RGB images, the addition of 100 spatially random depth samples reduces the prediction root-mean-square error by 50% on the NYU-Depth-v2 indoor dataset. It also boosts the percentage of reliable prediction from 59% to 92% on the KITTI dataset. We demonstrate two applications of the proposed algorithm: a plug-in module in SLAM to convert sparse maps to dense maps, and super-resolution for LiDARs. Software² and video demonstration³ are publicly available.



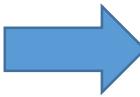
Sparse-to-Dense Propagation in Computer Vision

- Image labeling problems (segmentation)

**Sparse confident pixels
(user scribble)**

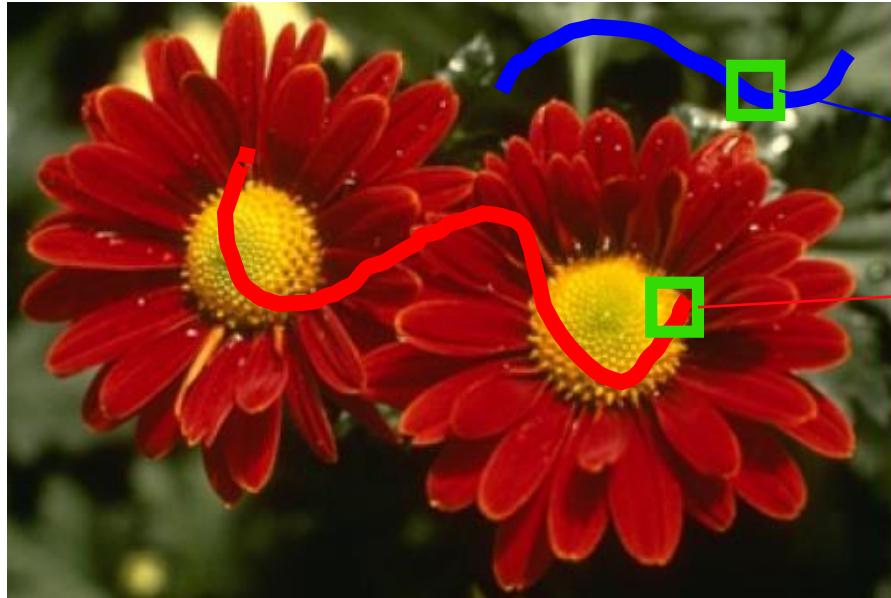


Dense segmentation result



Tutorial slides from Ľubor Ladický]

Example: Image Segmentation



How to formulate this?

Background : $x_i = 0$
Foreground : $x_i = 1$

Many algorithms are based on Graph Cuts.

Image Segmentation by Graph Cuts



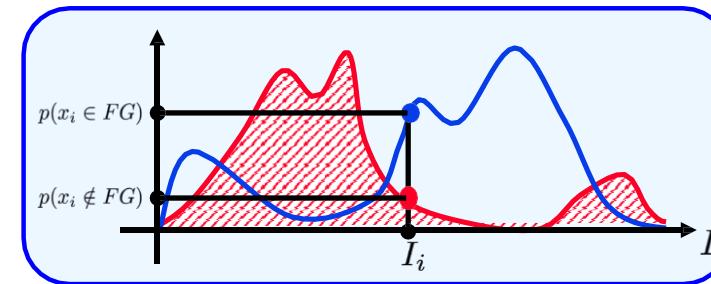
F/B based on data and color model

e.g. if color≈red or color≈yellow
 p is foreground
else
 p is background

Data term: based on the data (e.g. color) and model (e.g. color model)

Usually not perfect (noise, etc), so need **smoothness**

Should be estimated automatically or semi-automatically



Estimated using FG / BG color models

Image Segmentation by Graph Cuts



F/B based on neighbor labels

e.g. if neighbor of p is FG
 p is FG
else
 p is BG



Should get the same label

Smoothness term: based on the neighbor labels
i.e. “neighbouring pixels tend to take the same label”

Also high-level smoothness term or constraints

Image Segmentation by Graph Cuts

$$\mathbf{x}^* = \arg \min_{x \in \mathbf{L}} E(\mathbf{x})$$

$$E(\mathbf{x}) = \sum_{i \in \mathcal{V}} \psi_i(x_i) + \sum_{i \in \mathcal{V}, j \in \mathcal{N}_i} \psi_{ij}(x_i, x_j)$$

Data term **Smoothness term**

Unknowns

$x_i = 0 \implies i \in \text{Background}$

$x_i = 1 \implies i \in \text{Foreground}$

Other popular notation:

$l_i = 0 \implies i \in \text{Background}$

$l_i = 1 \implies i \in \text{Foreground}$

\mathcal{V} : the set of all the pixels

\mathcal{N}_i : the set of all the pixels in the neighborhood of pixel i
i.e. “the neighbors of i ”

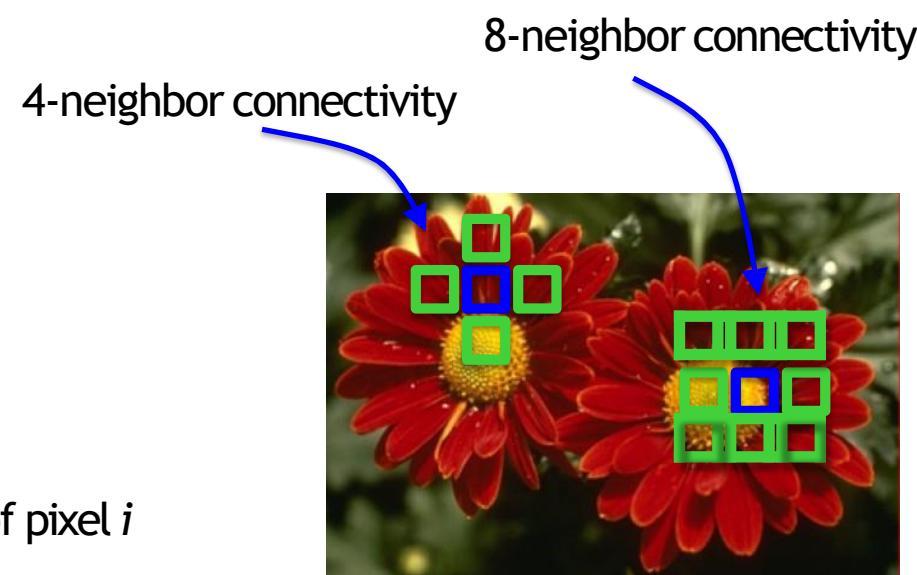


Image Segmentation by Graph Cuts

$$\mathbf{x}^* = \arg \min_{x \in \mathbf{L}} E(\mathbf{x})$$

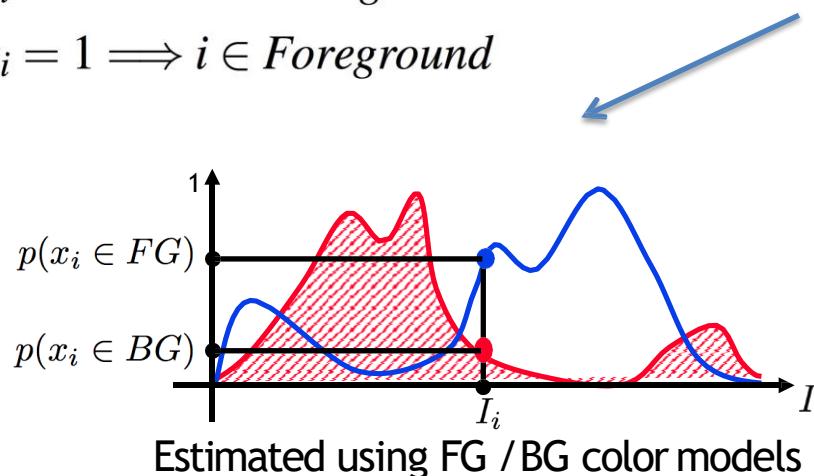
Unknowns

$x_i = 0 \implies i \in \text{Background}$

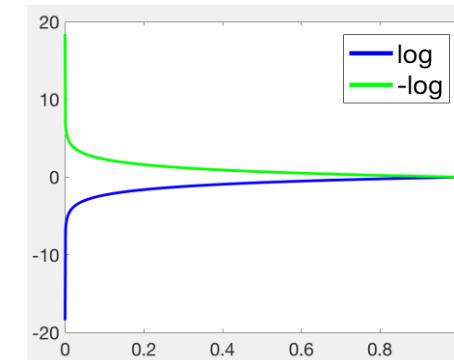
$x_i = 1 \implies i \in \text{Foreground}$

$$E(\mathbf{x}) = \sum_{i \in \mathcal{V}} \psi_i(x_i) + \sum_{i \in \mathcal{V}, j \in \mathcal{N}_i} \psi_{ij}(x_i, x_j)$$

Data term **Smoothness term**



$$\begin{aligned}\psi_i(0) &= \psi_i(BG) = -\log(p(x_i \in BG)) \\ \psi_i(1) &= \psi_i(FG) = -\log(p(x_i \in FG))\end{aligned}$$



Why -log?

- we are adding weights. We multiply probabilities, so add logs
- maximize probabilities, so minimize -log

Image Segmentation by Graph Cuts

$$\mathbf{x}^* = \arg \min_{x \in \mathbf{L}} E(\mathbf{x})$$

Unknowns

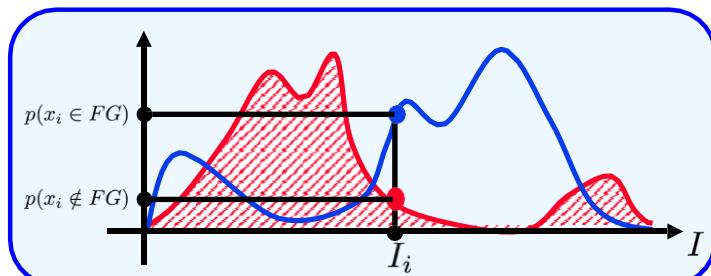
$x_i = 0 \implies i \in \text{Background}$

$x_i = 1 \implies i \in \text{Foreground}$

$$E(\mathbf{x}) = \sum_{i \in \mathcal{V}} \psi_i(x_i) + \sum_{i \in \mathcal{V}, j \in \mathcal{N}_i} \psi_{ij}(x_i, x_j)$$

Data term **Smoothness term**

$$\begin{aligned}\psi_i(0) &= -\log(p(x_i \notin FG)) \\ \psi_i(1) &= -\log(p(x_i \in FG))\end{aligned}$$



Estimated using FG / BG color models

$$\psi_{ij}(x_i, x_j) = \delta(x_i \neq x_j) \quad \begin{array}{l} \text{if } x_i = x_j \Rightarrow 0 \\ \text{else } 1 \end{array}$$

$$\psi_{ij}(x_i, x_j) = K_{ij} \delta(x_i \neq x_j)$$

where $K_{ij} = \lambda_1 + \lambda_2 \exp(-\beta(I_i - I_j)^2)$

Intensity dependent smoothness

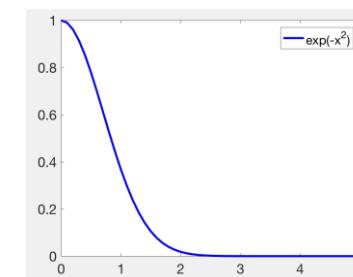


Image Segmentation by Graph Cuts

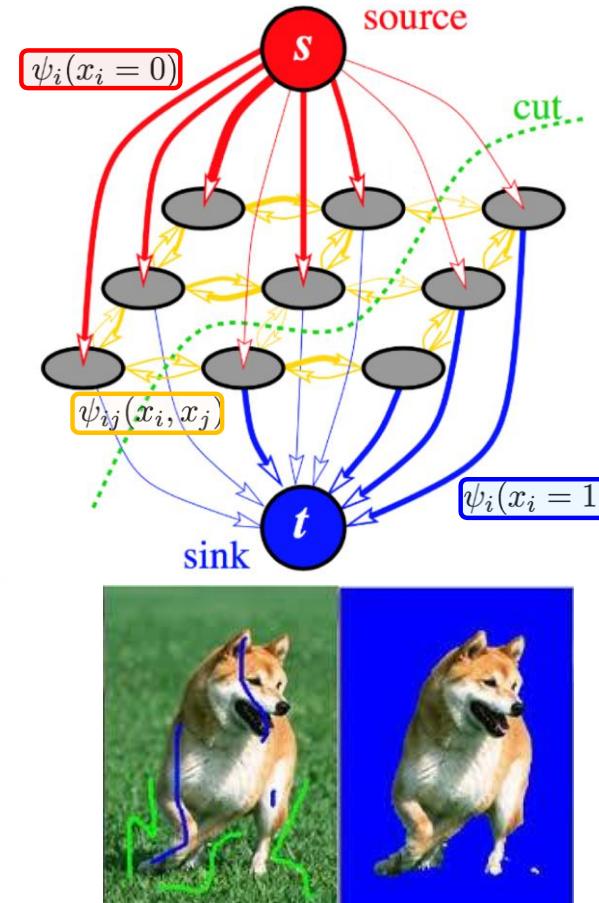
$$E(\mathbf{x}) = \sum_{i \in \mathcal{V}} \psi_i(x_i) + \sum_{i \in \mathcal{V}, j \in \mathcal{N}_i} \psi_{ij}(x_i, x_j)$$

Data term **Smoothness term**

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \mathbf{L}} E(\mathbf{x})$$

How to solve this optimization problem?

- Transform into max-flow / min-cut problem
- Solve it using max-flow / min-cut algorithm



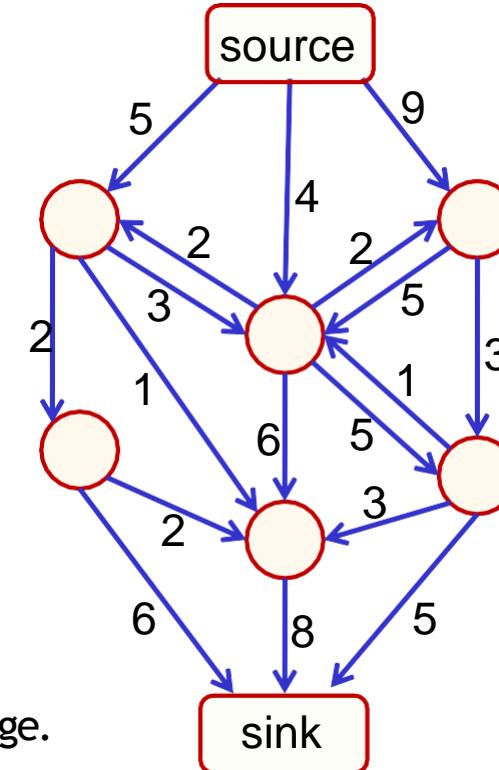
Max Flow Problem

Task :

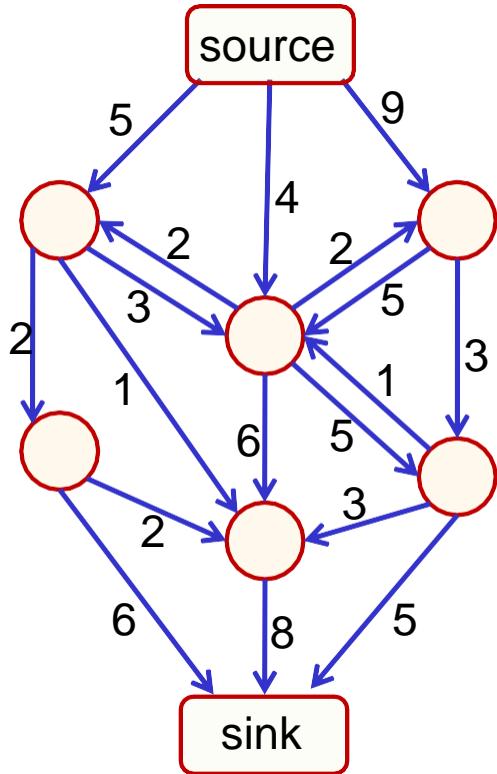
Maximize the flow from the source to the sink such that

- 1) The flow is conserved for each node ($\text{in}=\text{out}$)
- 2) The flow for each pipe does not exceed the capacity

“capacity” of an edge: the maximum amount of flow that can pass through an edge.



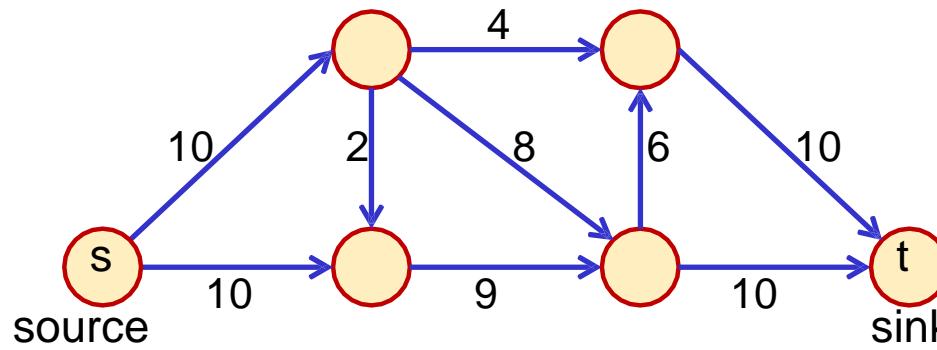
Max Flow Problem



$$\begin{aligned} & \text{flow from the source} \\ & \text{flow from node } i \text{ to node } j \\ & \max \sum_{i \in V} f_{si} \\ \text{s.t.} \quad & 0 \leq f_{ij} \leq c_{ij}, \\ & \sum_{j \in N(i)} f_{ji} - f_{ij} = 0, \\ & \forall (i, j) \in E \\ & \forall i \in V \setminus \{s, t\} \\ & \text{conservation of flow} \\ & \text{i.e. flow in} = \text{flow out} \end{aligned}$$

Max Flow Problem

Graph with capacity

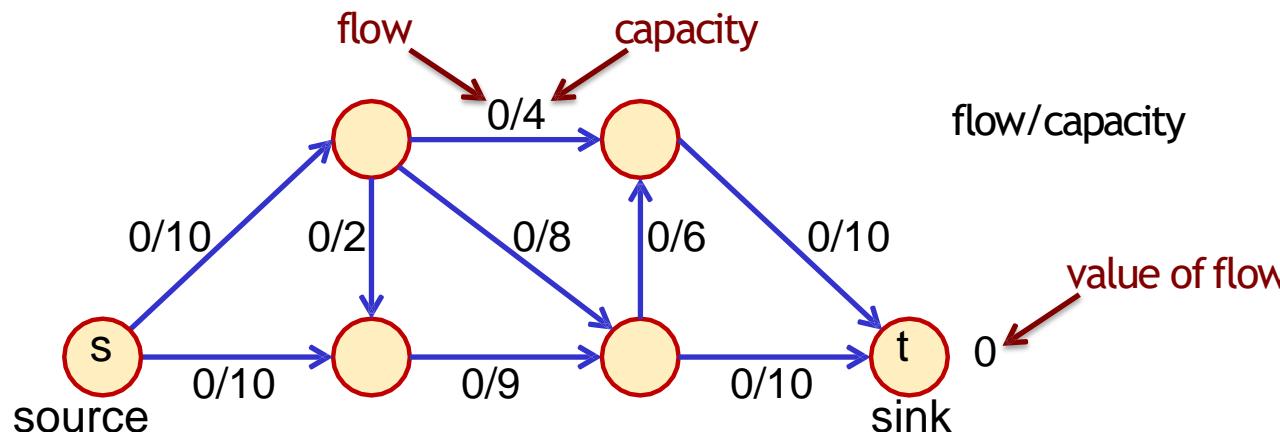


flow

capacity

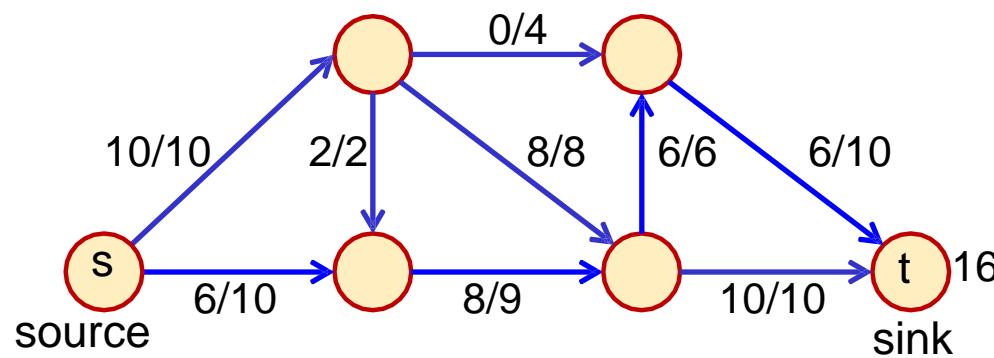
flow/capacity

value of flow



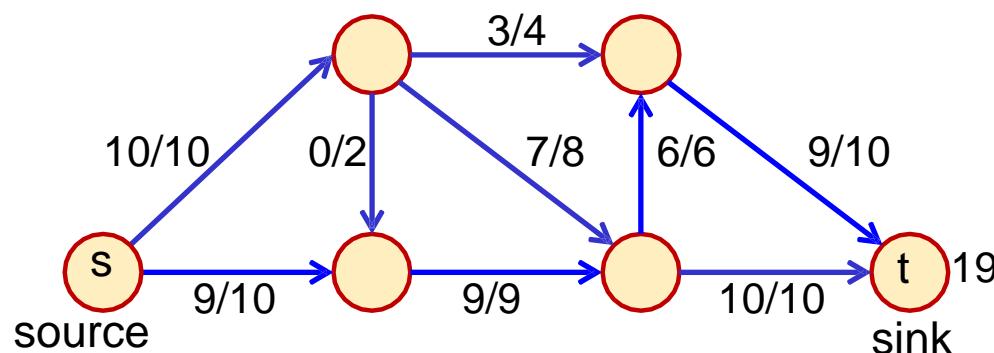
Max Flow Problem

flow/capacity



flow in=flow out
flow conversation
flow<=capacity

Total flow=16



Total flow=19

Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

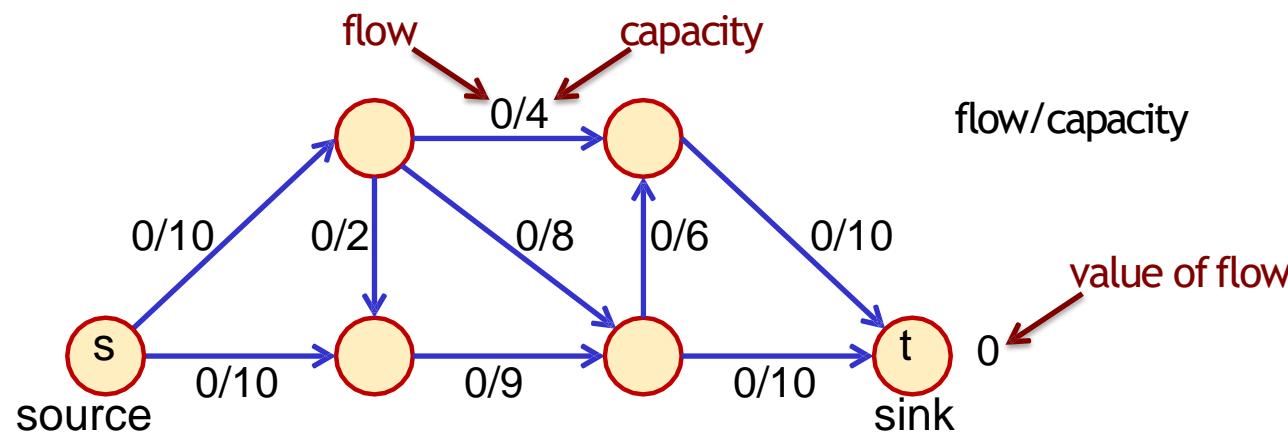
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

End



<http://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/07NetworkFlowI-2x2.pdf>

Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

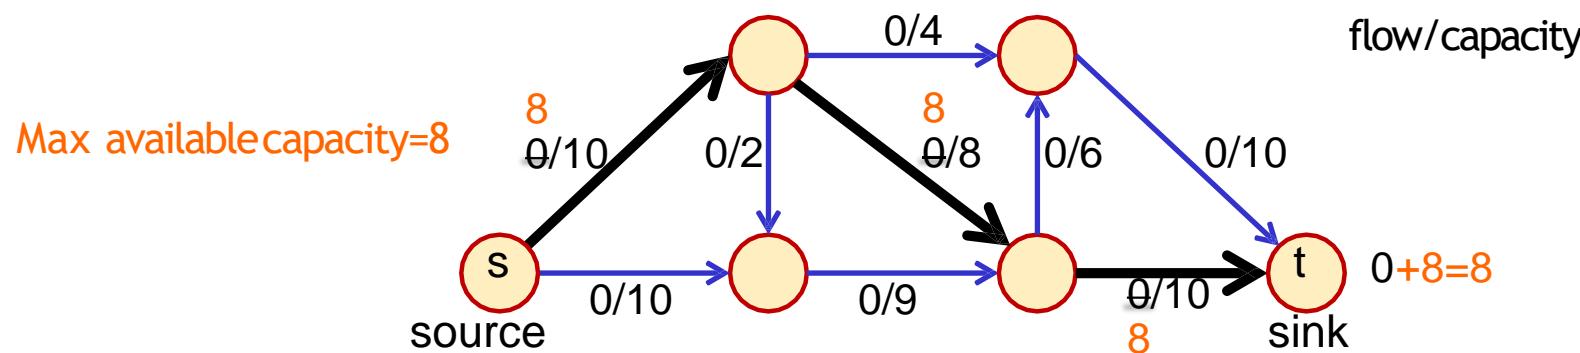
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

End



<http://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/07NetworkFlowI-2x2.pdf>

Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

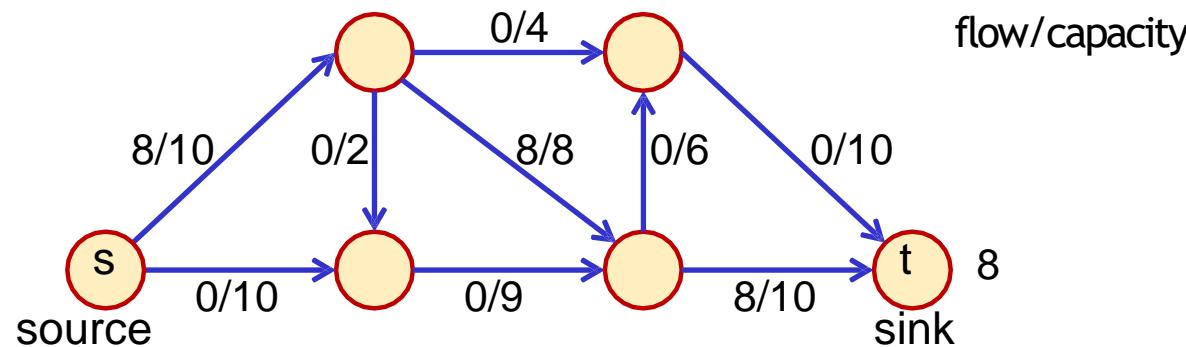
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

End



<http://www.cs.princeton.edu/~wayne/kleinberg-tardos/pdf/07NetworkFlowI-2x2.pdf>

Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

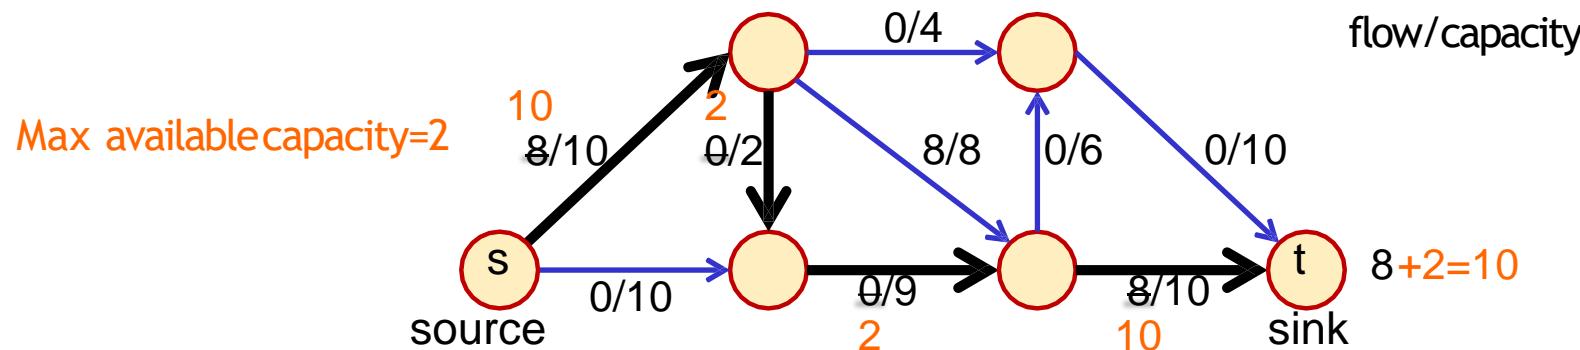
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

End



Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

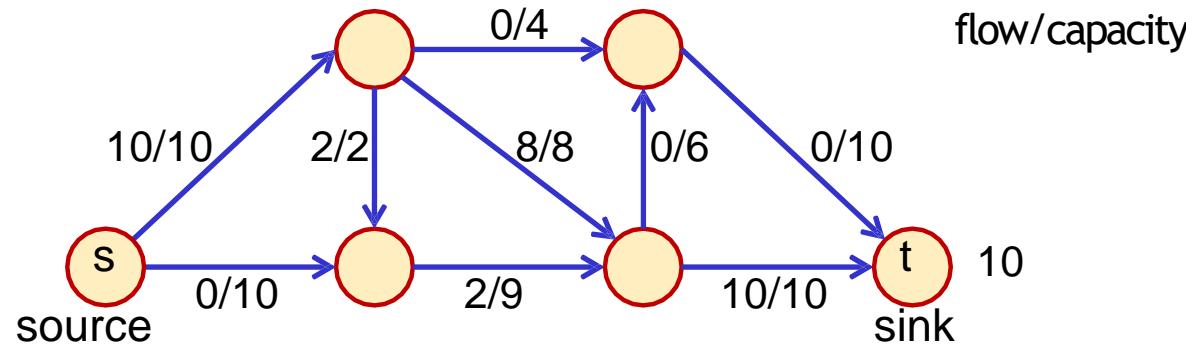
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

End



Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

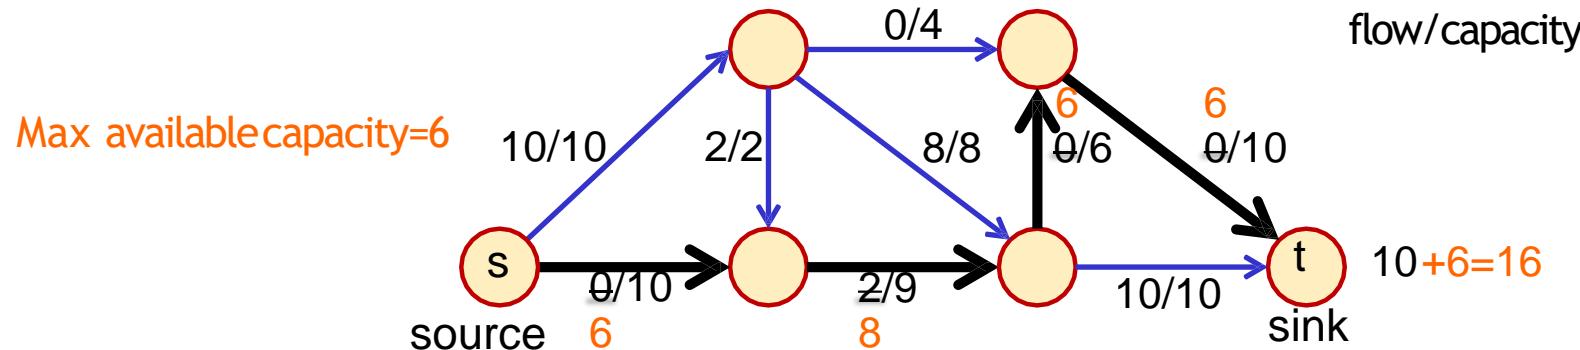
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

End



Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

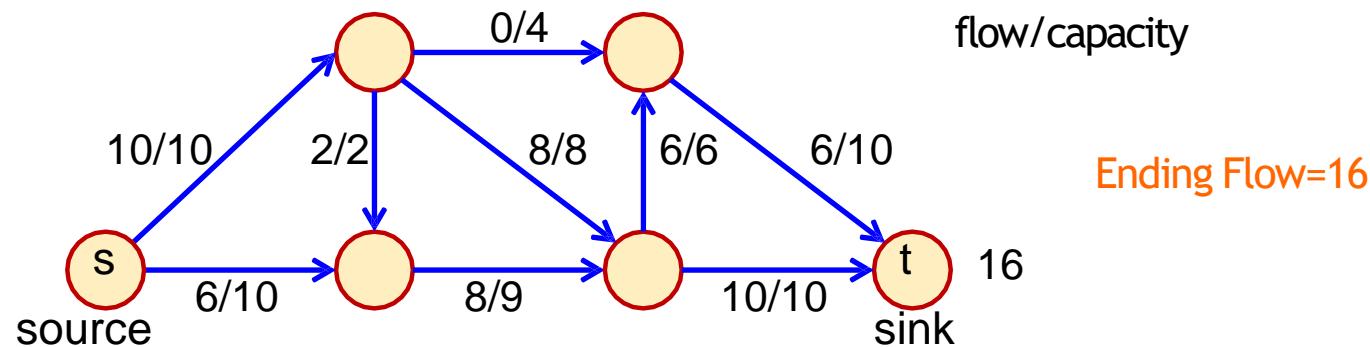
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

End



Max Flow Problem - Greedy algorithm

Greedy algorithm:

flow=0

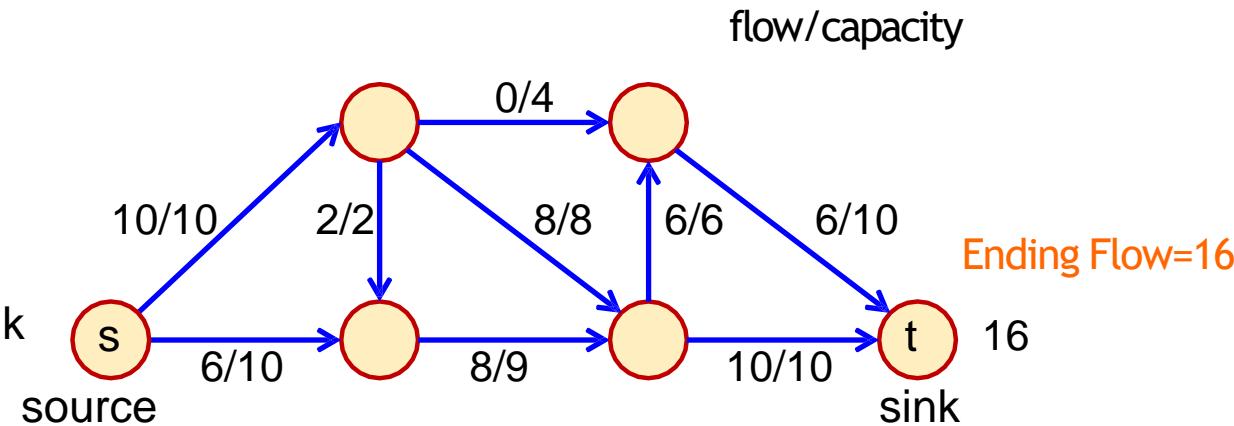
Find a “valid” path from source to sink

While (path exists)

“Augment” flow

Find a “valid” path from source to sink

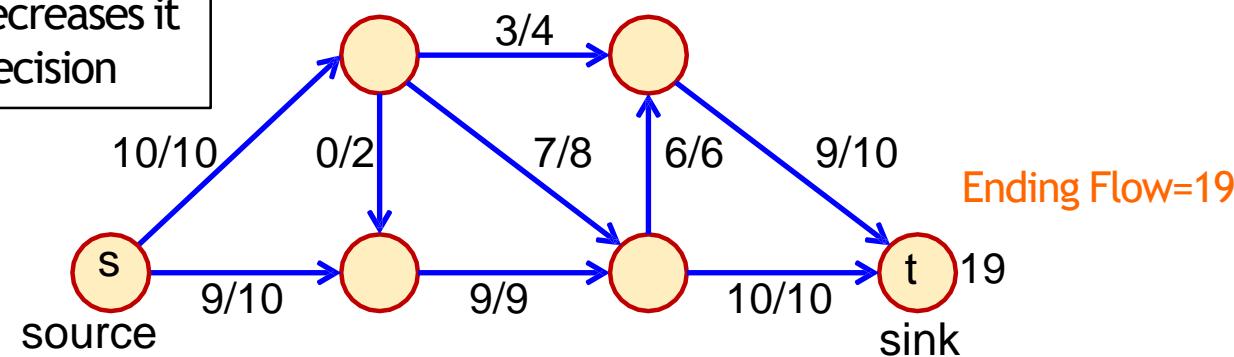
End



Why does it fail?

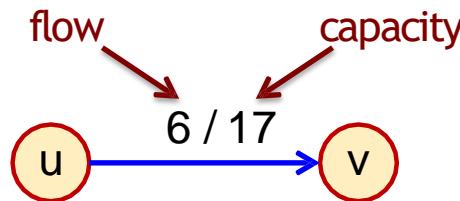
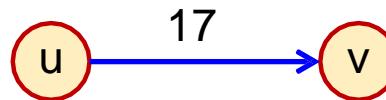
→ Once it increases flow on an edge, it never decreases it

→ We need some mechanisms to “undo” bad decision

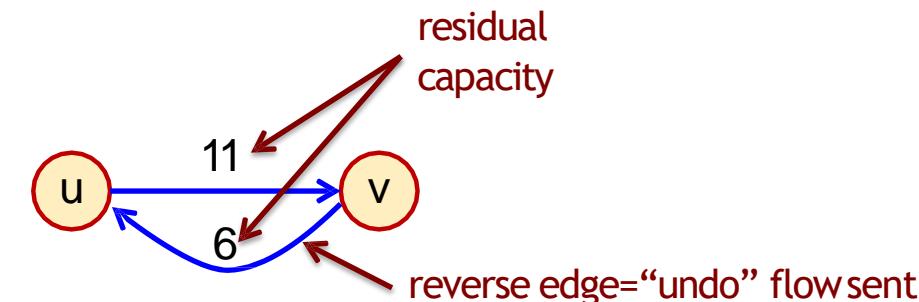
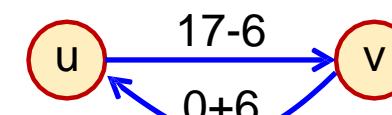
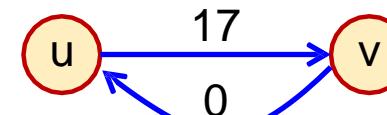


Max Flow Problem – Residual Network

Original network G :



Residual network G_f :



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

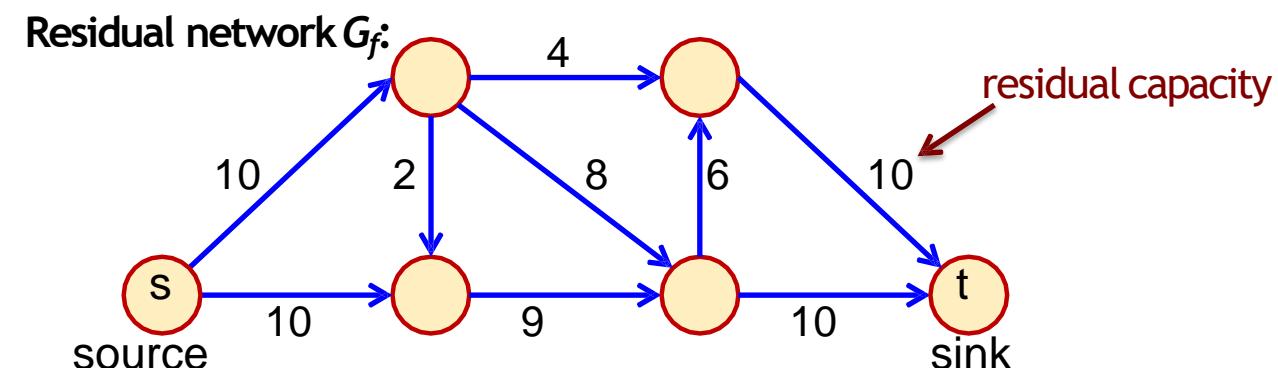
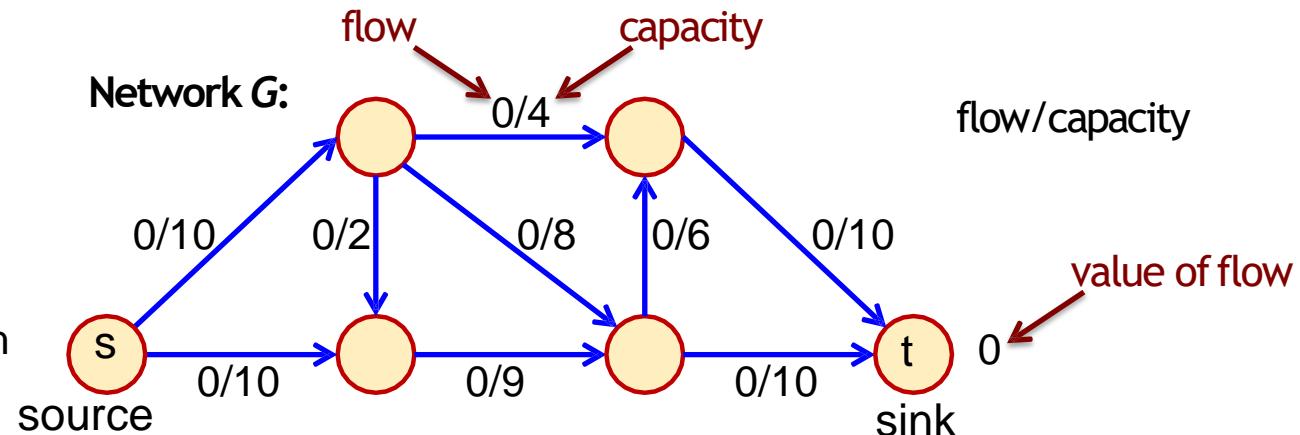
While (path exists)

 flow += maximum capacity in the path

 Build the residual graph

 Find a path in the residual graph

End



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

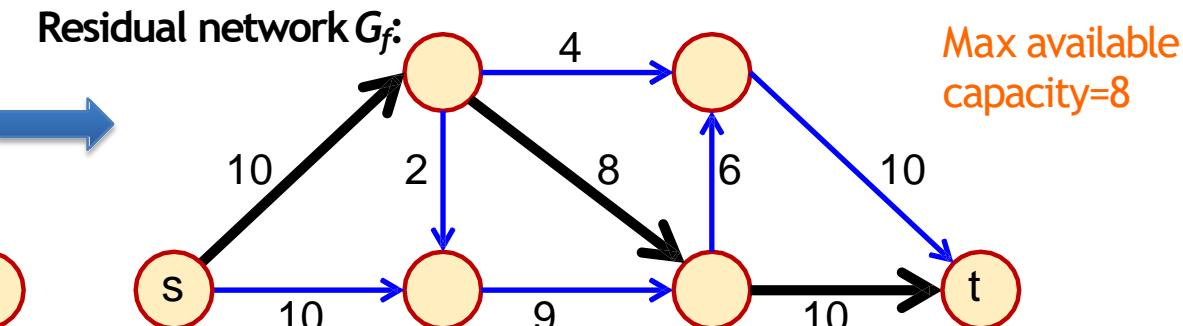
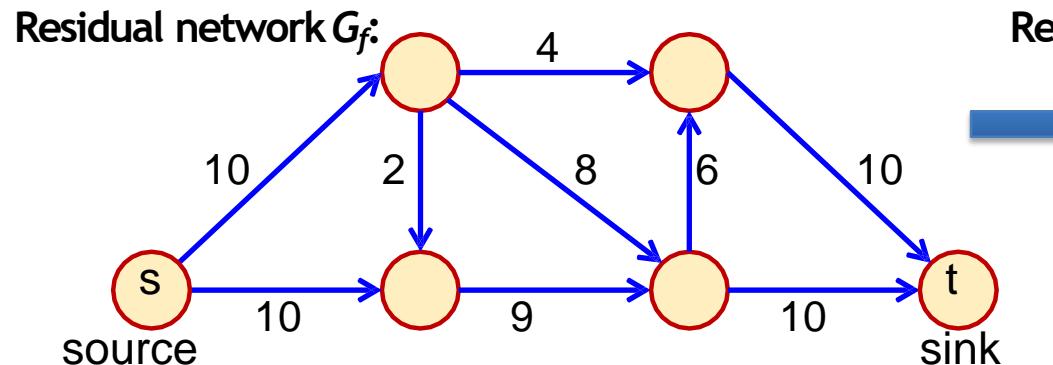
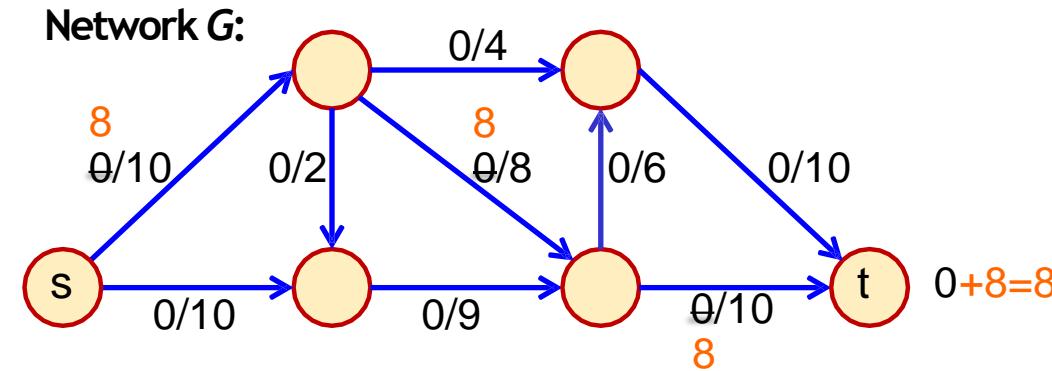
While (path exists)

 flow += maximum capacity in the path

 Build the residual graph

 Find a path in the residual graph

End



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

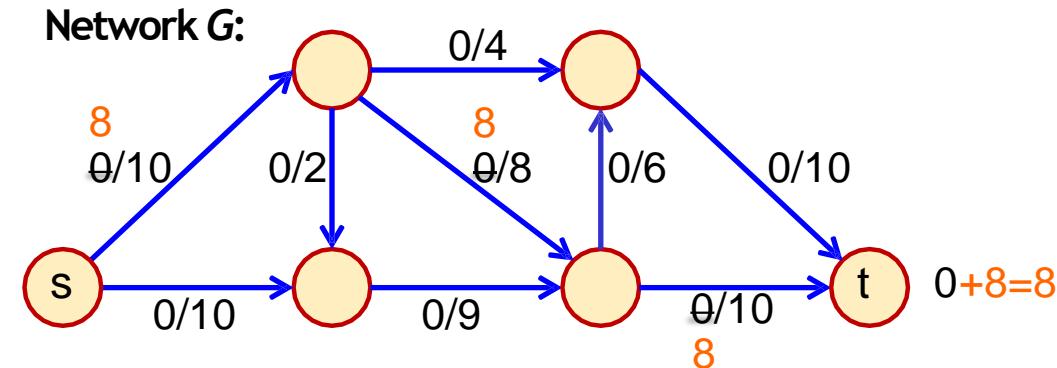
While (path exists)

 flow += maximum capacity in the path

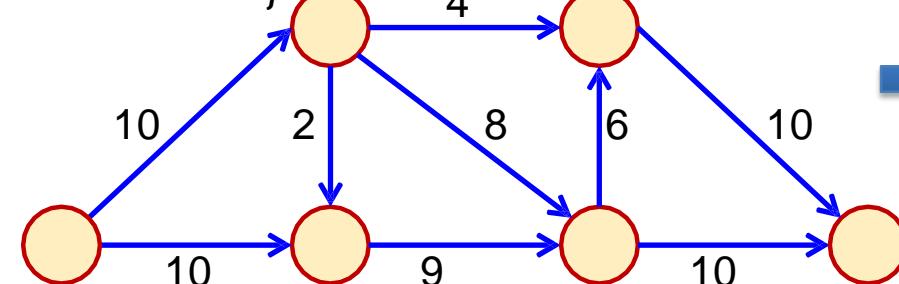
 Build the residual graph

 Find a path in the residual graph

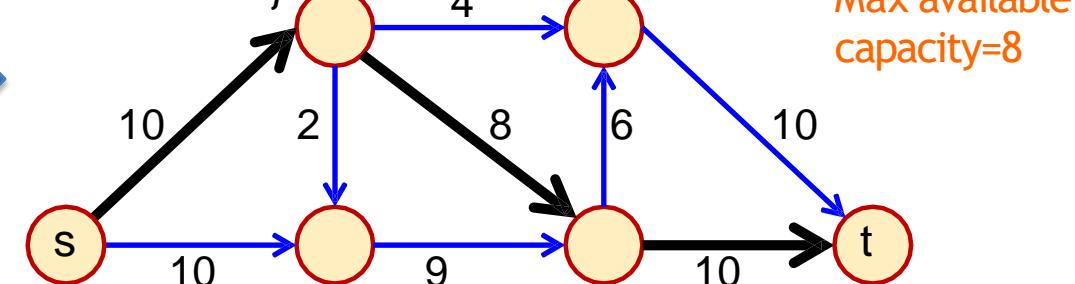
End



Residual network G_f :

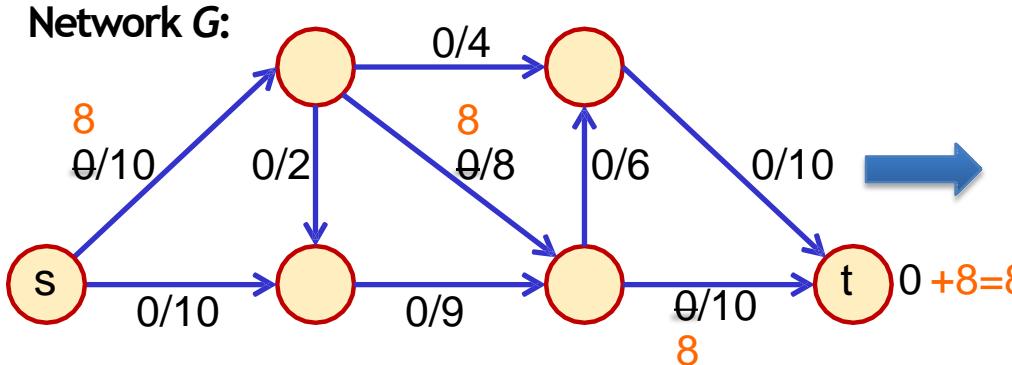


Residual network G_f :

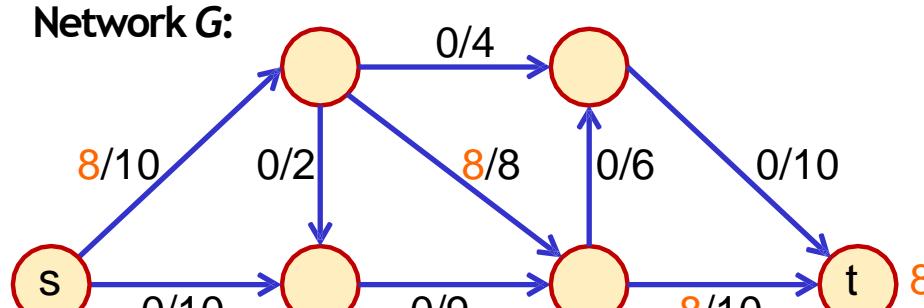


Max Flow Problem – Ford & Fulkerson

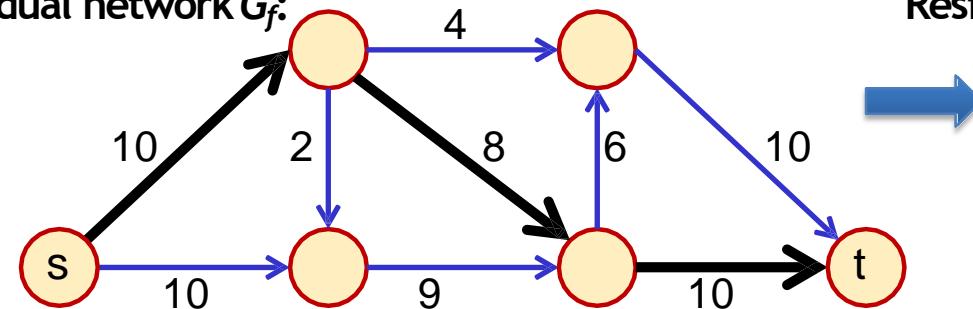
flow/capacity



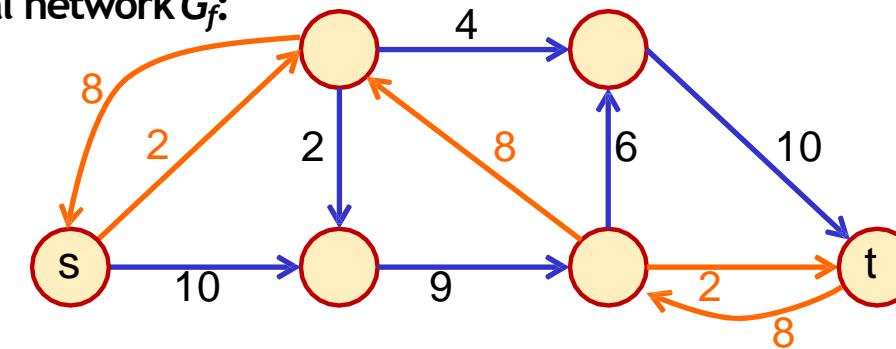
Max available capacity=8



Residual network G_f :



Residual network G_f :



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

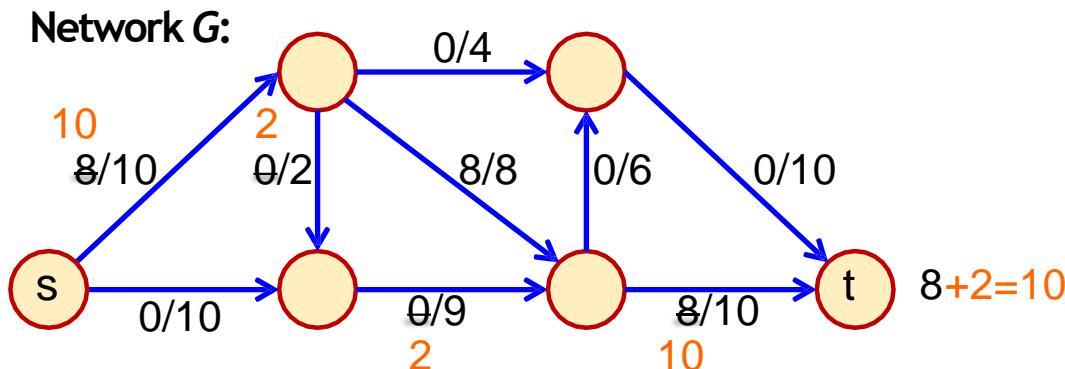
While (path exists)

 flow += maximum capacity in the path

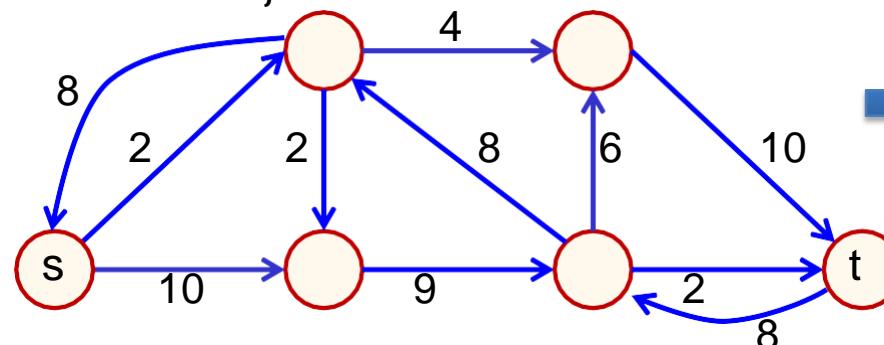
 Build the residual graph

 Find a path in the residual graph

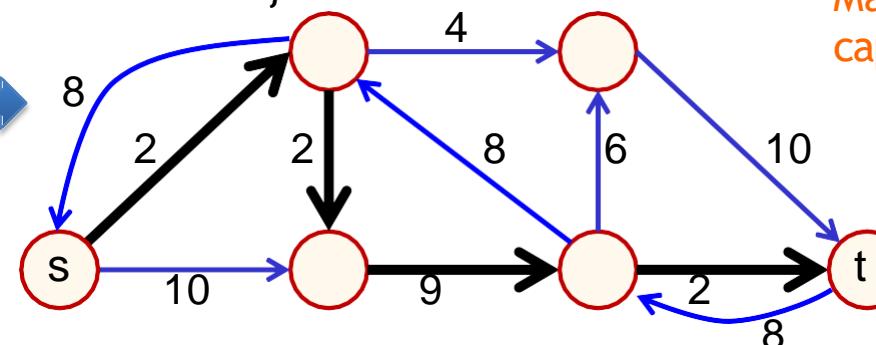
End



Residual network G_f :



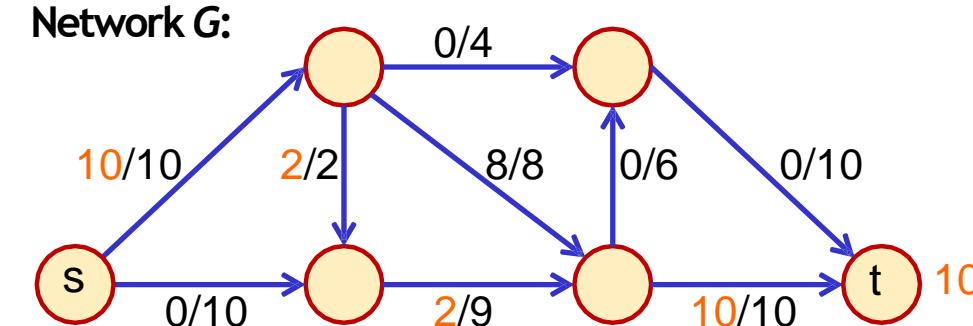
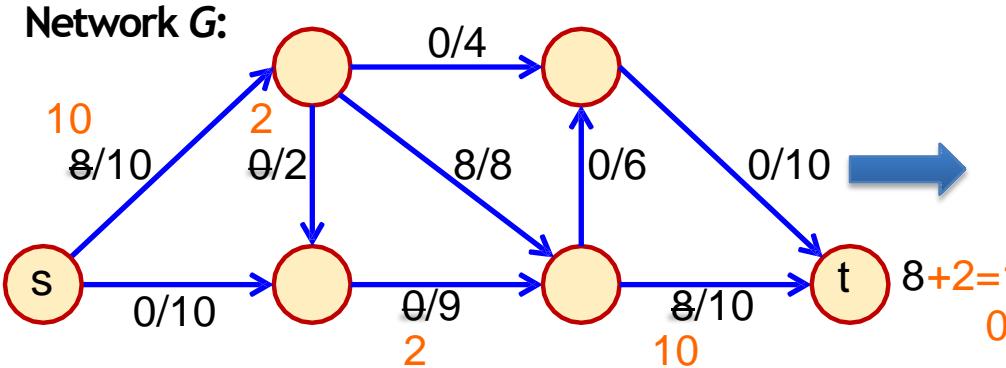
Residual network G_f :



Max available capacity=2

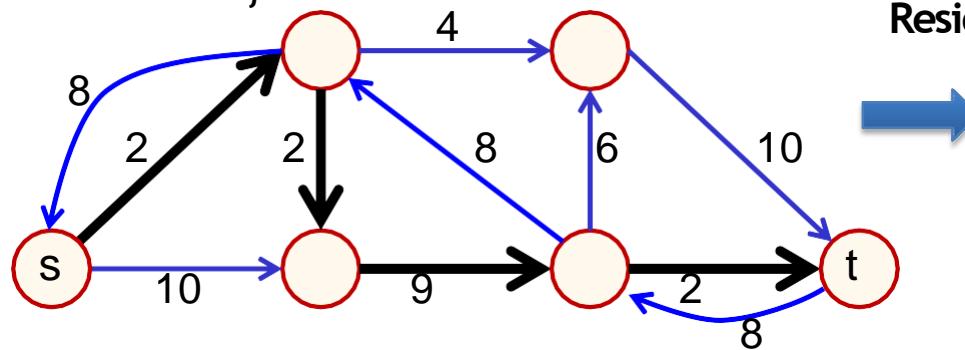
Max Flow Problem – Ford & Fulkerson

flow/capacity

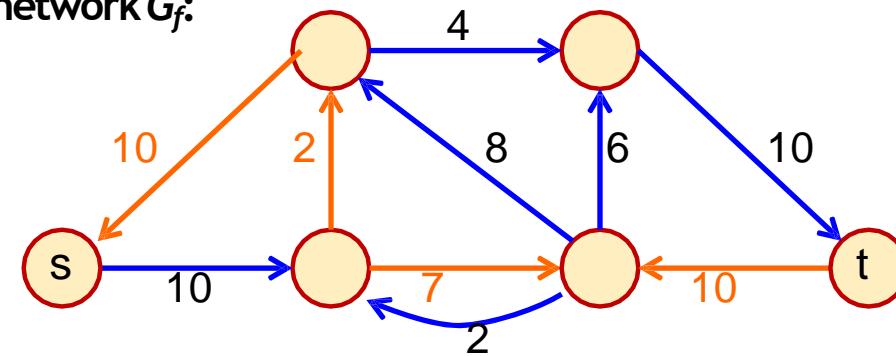


Max available capacity=2

Residual network G_f :



Residual network G_f :



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

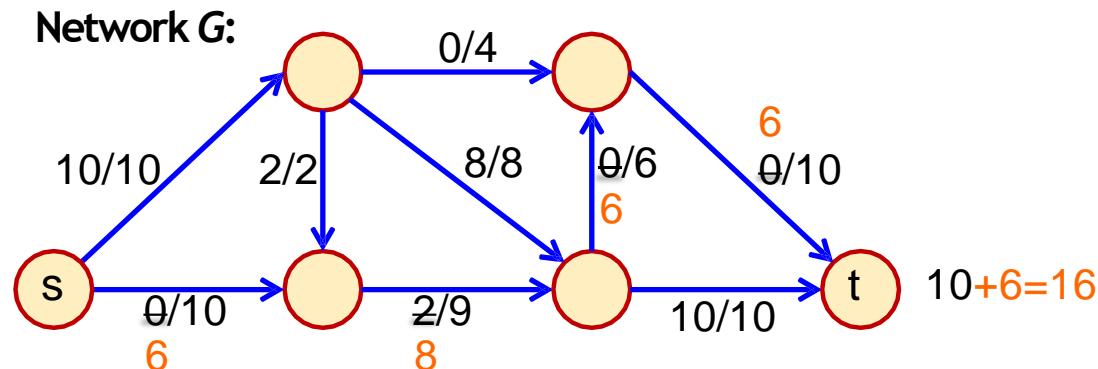
While (path exists)

 flow += maximum capacity in the path

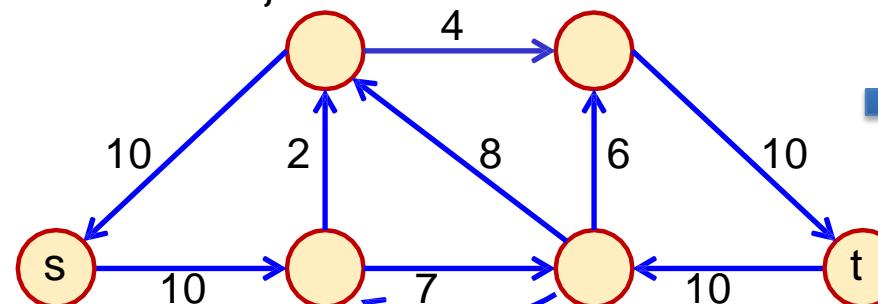
 Build the residual graph

 Find a path in the residual graph

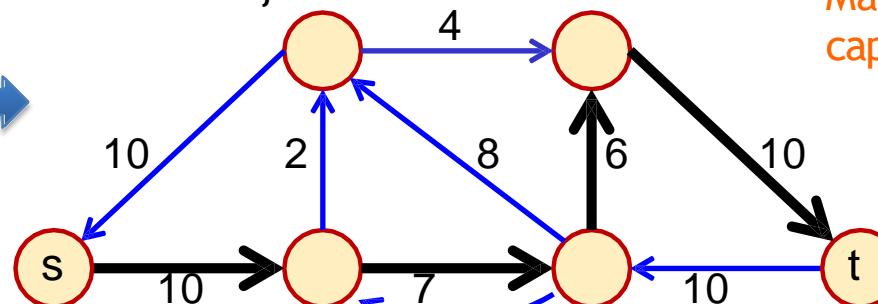
End



Residual network G_f :



Residual network G_f :



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

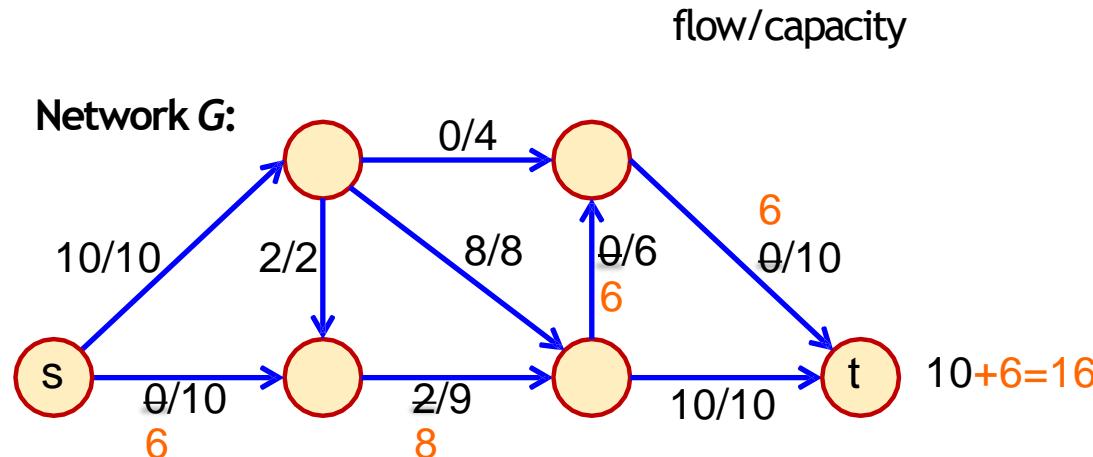
While (path exists)

 flow += maximum capacity in the path

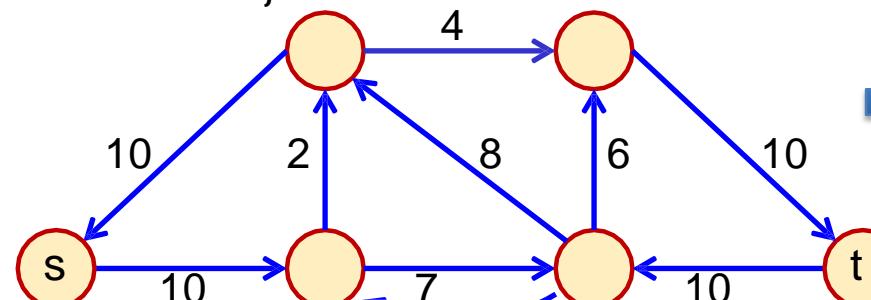
 Build the residual graph

 Find a path in the residual graph

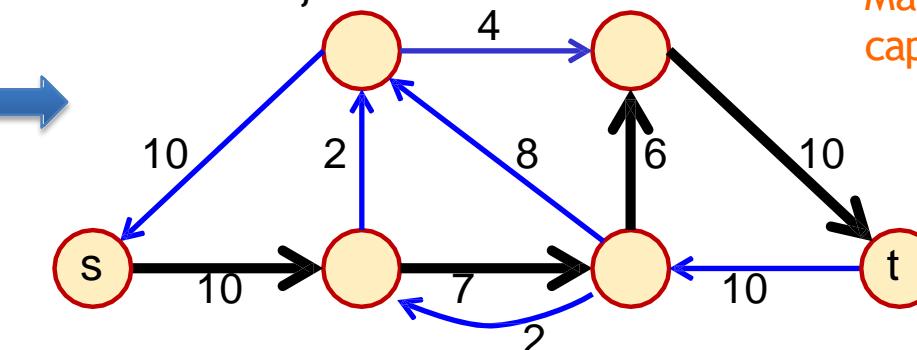
End



Residual network G_f :

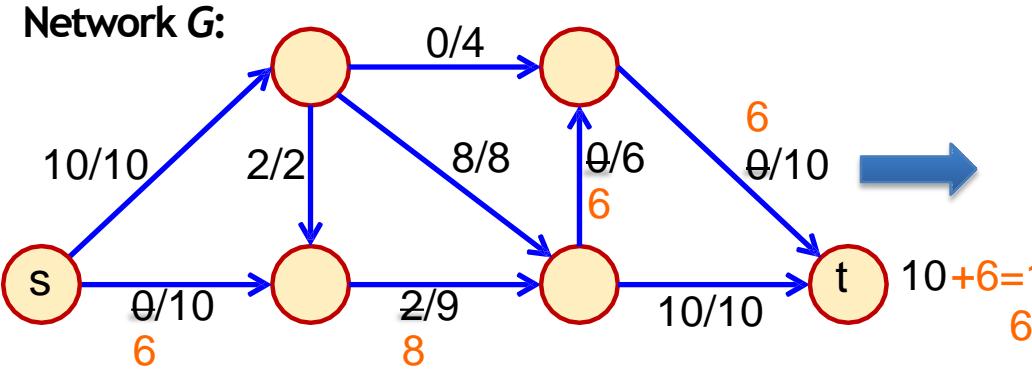


Residual network G_f :



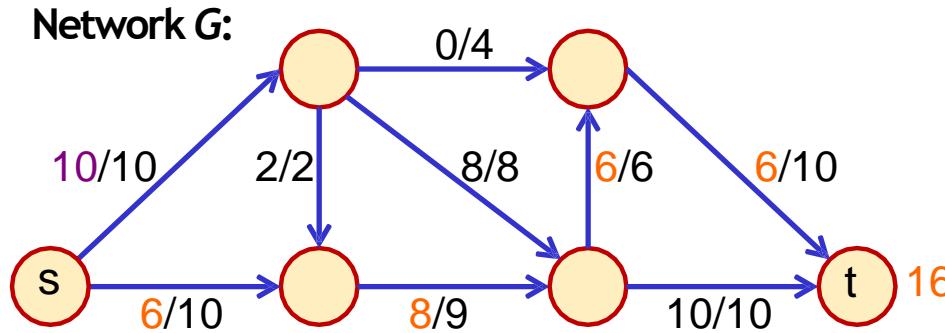
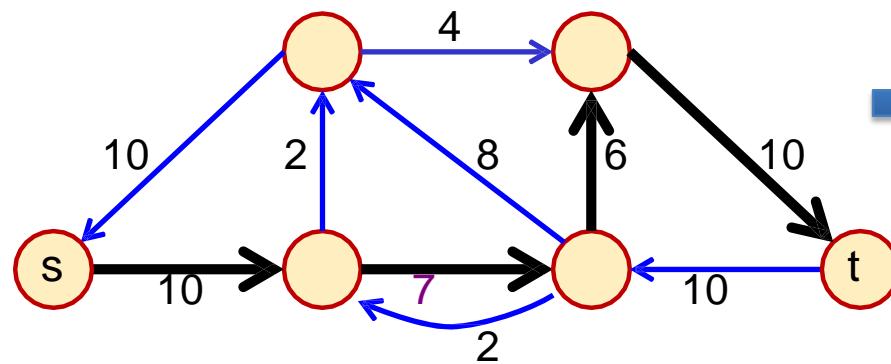
Max Flow Problem – Ford & Fulkerson

flow/capacity

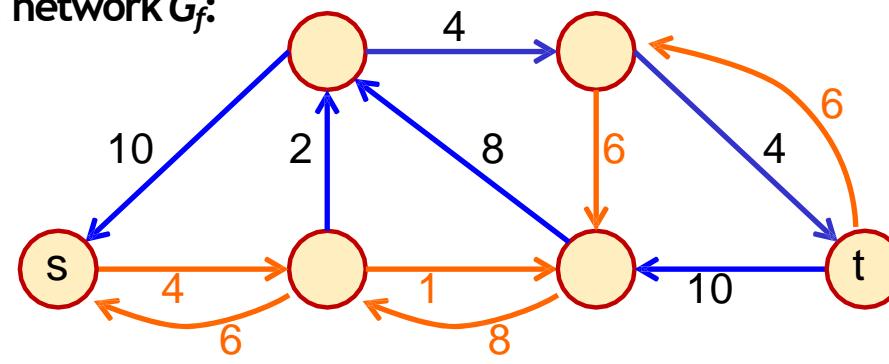


Max available capacity=6

Residual network G_f :



Residual network G_f :



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

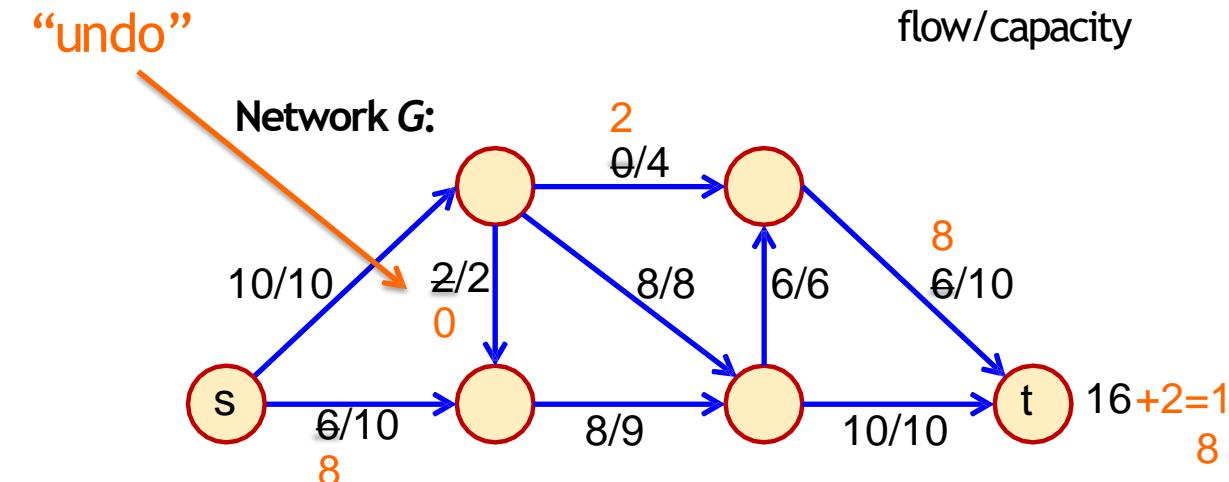
While (path exists)

 flow += maximum capacity in the path

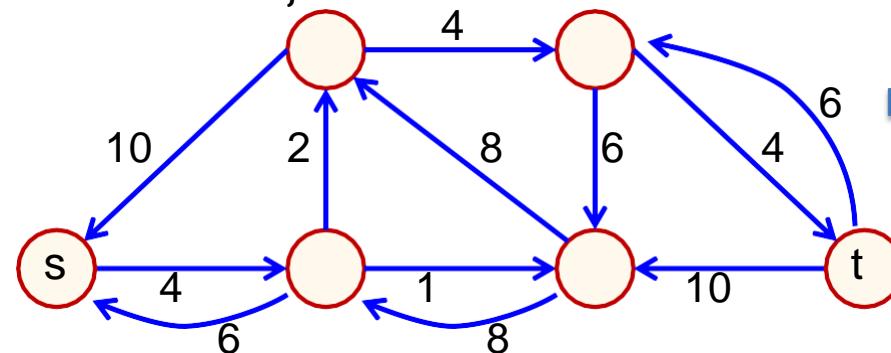
 Build the residual graph

 Find a path in the residual graph

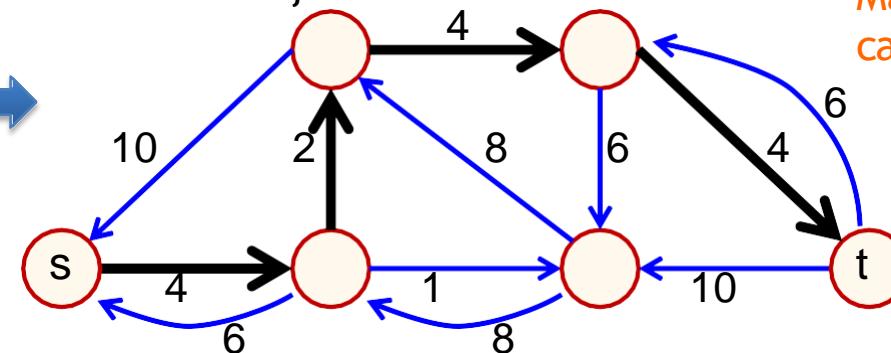
End



Residual network G_f :

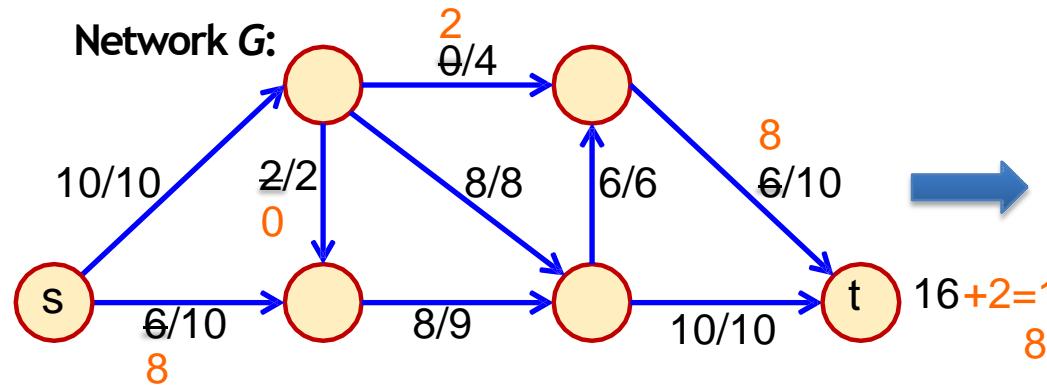


Residual network G_f :



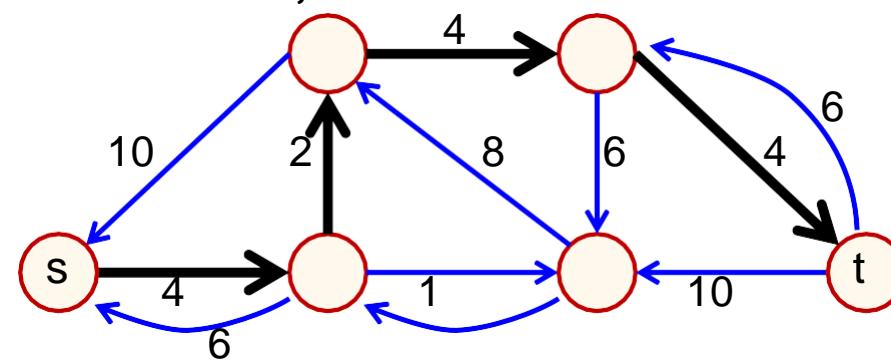
Max Flow Problem – Ford & Fulkerson

flow/capacity

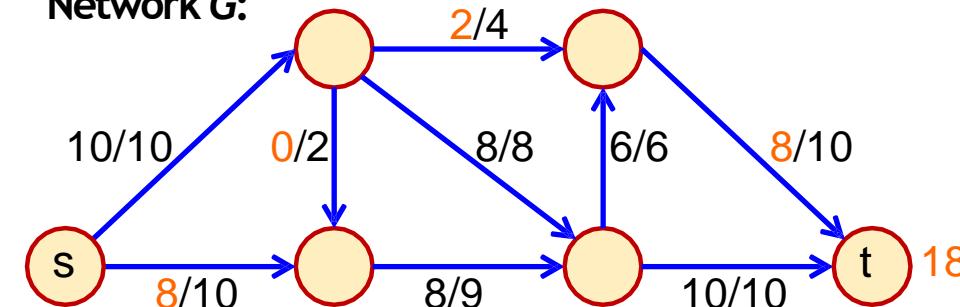


Max available capacity = 2

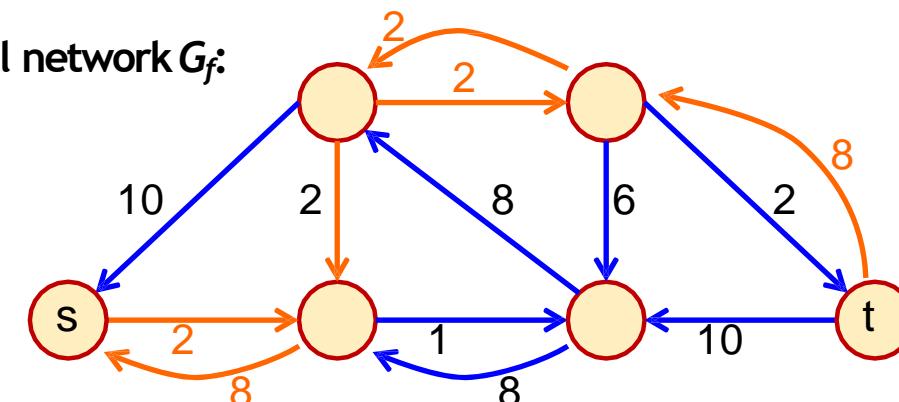
Residual network G_f :



Network G:



Residual network G_f :



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

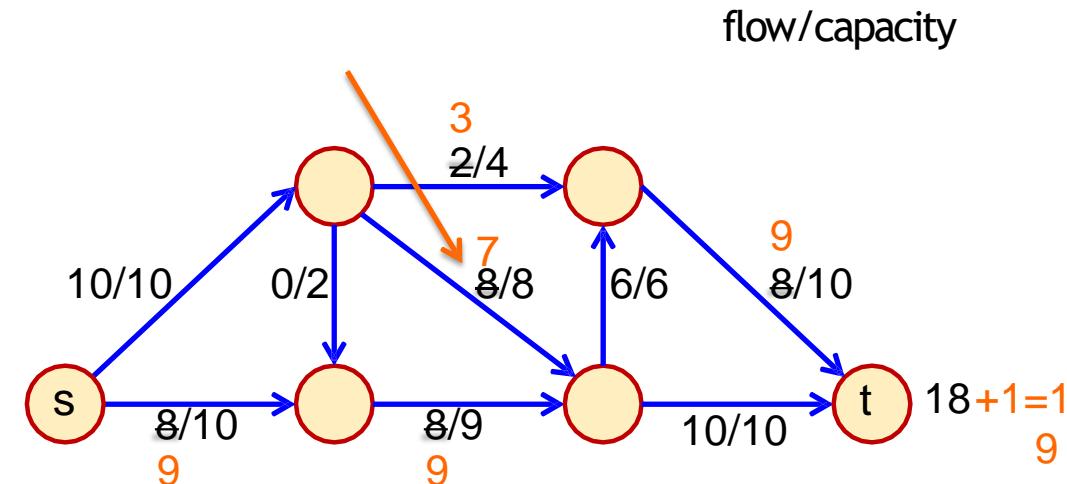
While (path exists)

 flow += maximum capacity in the path

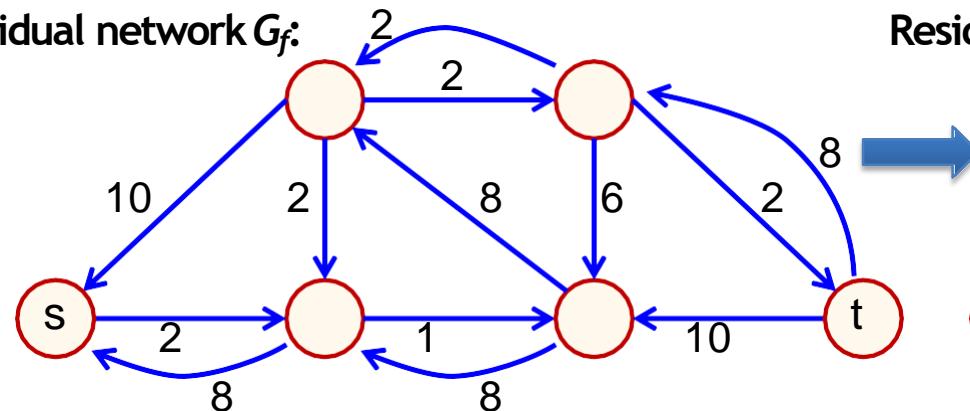
 Build the residual graph

 Find a path in the residual graph

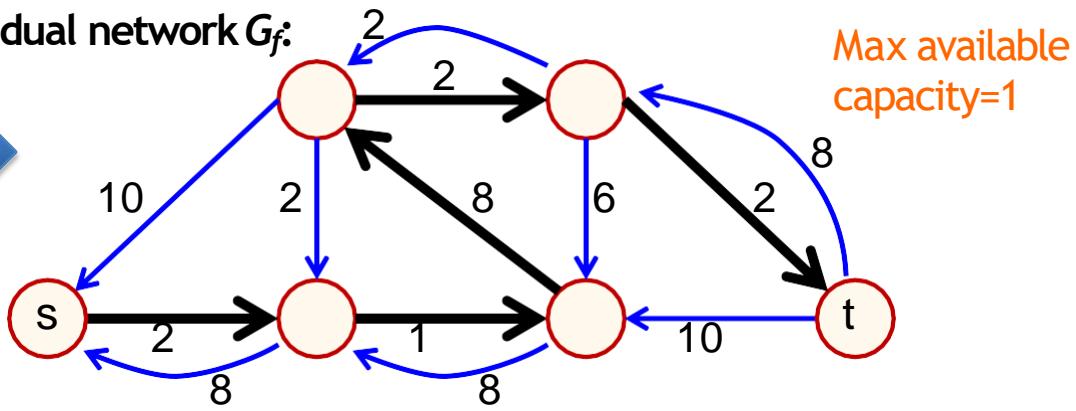
End



Residual network G_f :

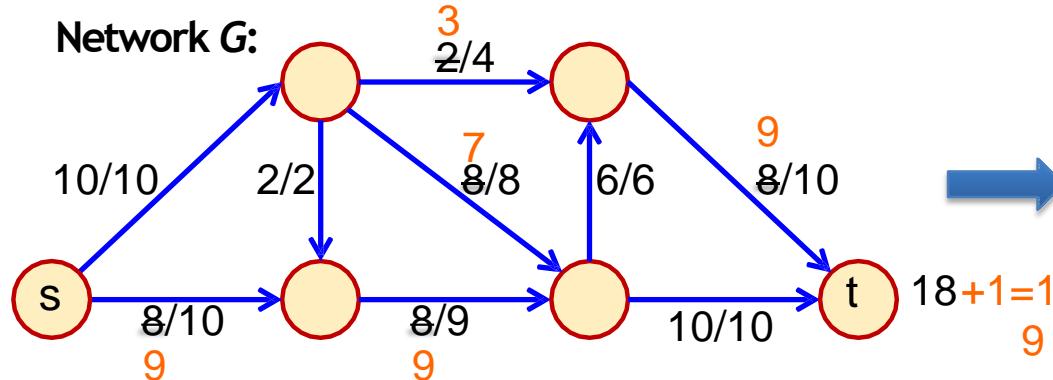


Residual network G_f :

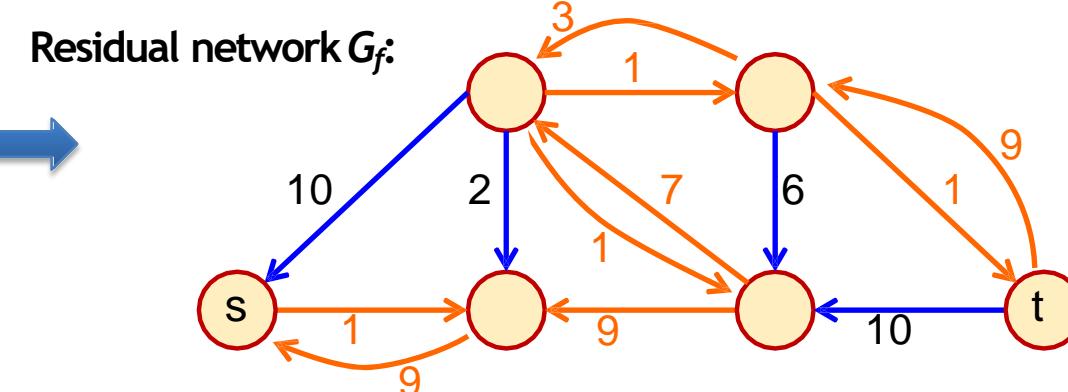
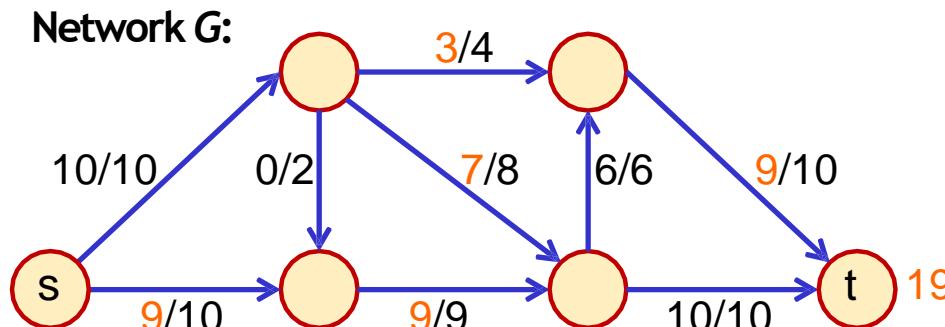
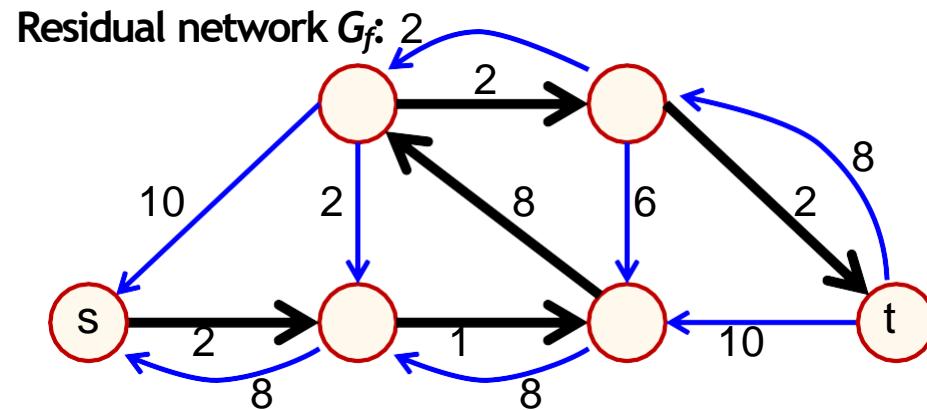


Max Flow Problem – Ford & Fulkerson

flow/capacity



Max available capacity=1



Max Flow Problem – Ford & Fulkerson

Ford & Fulkerson algorithm (1956)

flow=0

Find a path from source to sink

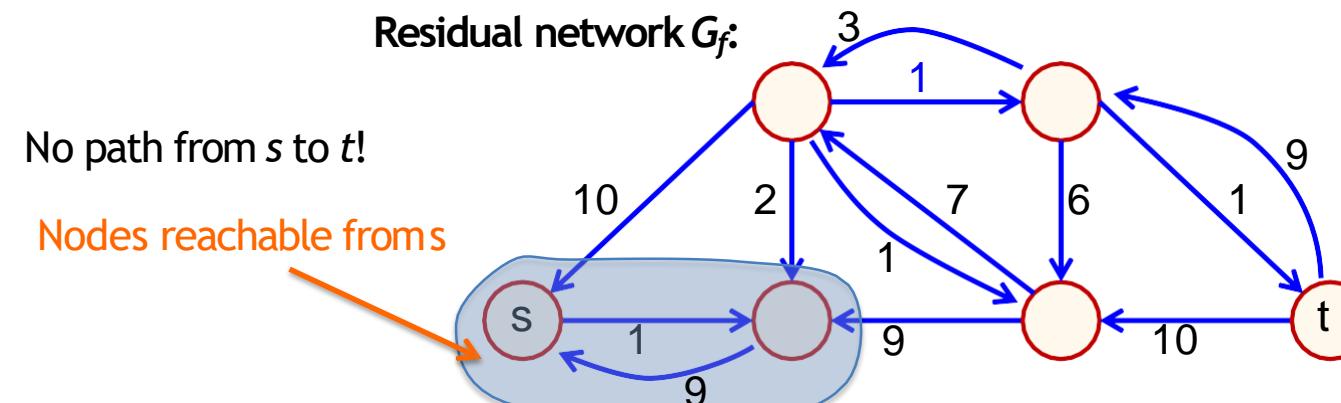
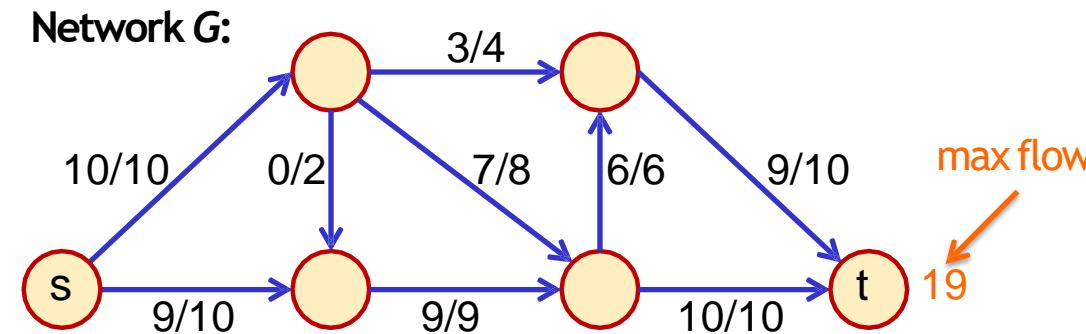
While (path exists)

 flow += maximum capacity in the path

 Build the residual graph

 Find a path in the residual graph

End



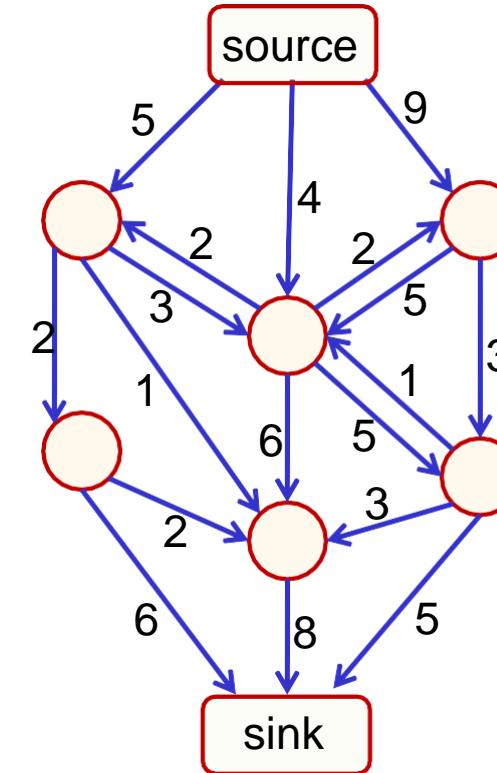
Min Cut Problem – Image Segmentation

* Min-Cut

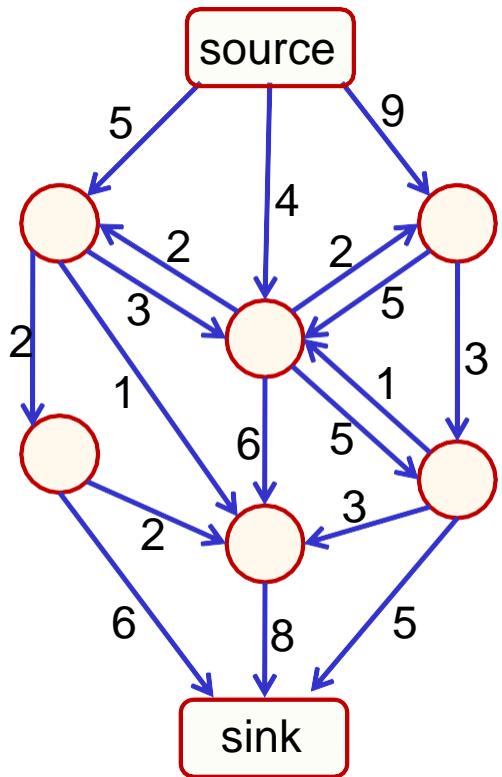
Task :

Minimize the cost of the cut

- 1) Each node is either assigned to the source S or sink T
- 2) The cost of the edge (i, j) is taken if $(i \in S)$ and $(j \in T)$



Min Cut Problem – Image Segmentation



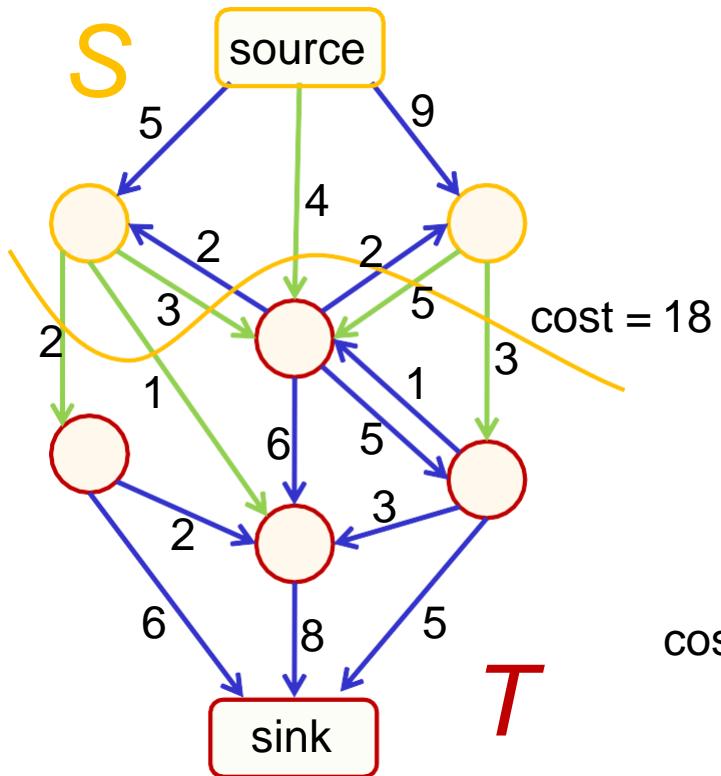
$$\min_{S,T} \sum_{i \in S, j \in T} c_{ij}$$

edge costs

source set sink set

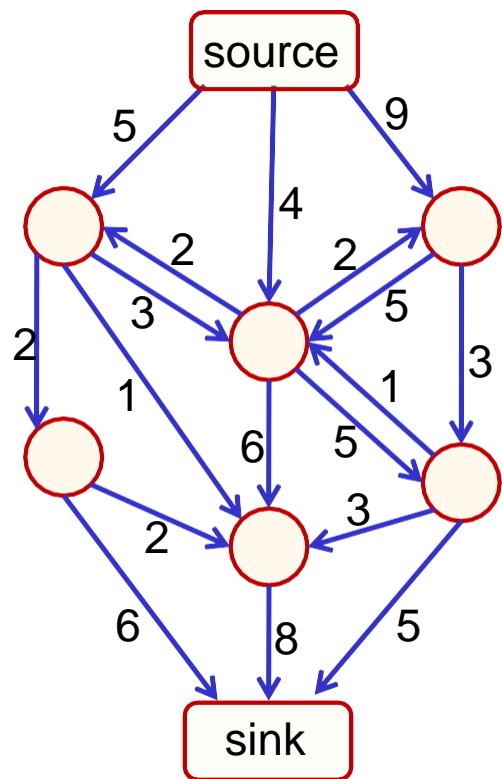
s.t. $s \in S, t \in T$

Min Cut Problem – Image Segmentation



$$\begin{array}{ll}
 \text{edge costs} & \\
 \min_{S,T} \sum_{i \in S} \sum_{j \in T} c_{ij} & s.t. \quad s \in S, \quad t \in T \\
 \text{source set} & \text{sink set}
 \end{array}$$

Min Cut Problem – Image Segmentation



Find the min-cut
i.e. find the partition that provides
the lowest cut cost

$$x_i = 0 \implies x_i \in S$$

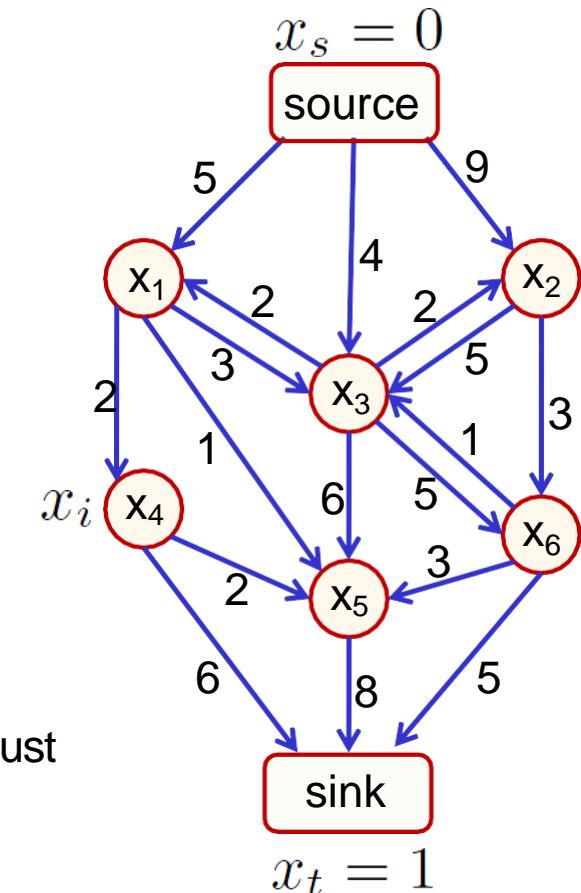
$$x_i = 1 \implies x_i \in T$$

Find the binary “partition variables”

* **Min-cut solution is equivalent to Max-Flow solution.**

→ A cut is a bottleneck; any flow through the network must pass all paths through the bottleneck at once.

→ The total flow is limited by the narrow part, minimum cut.

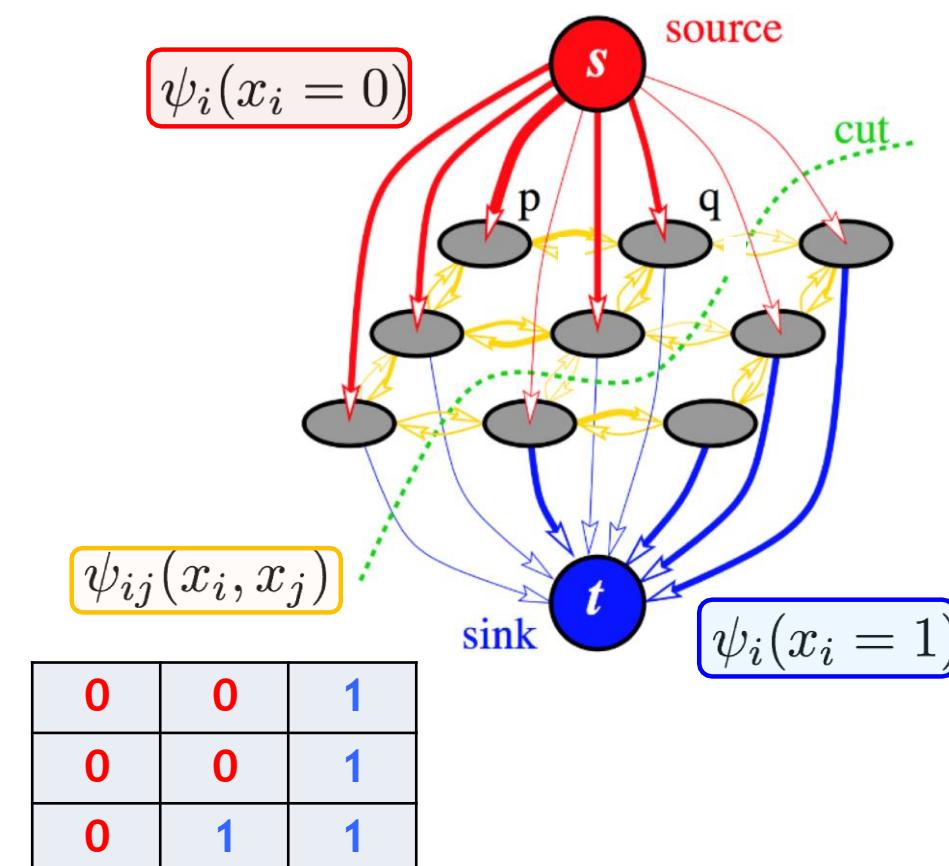


Min Cut Problem – Image Segmentation

$$E(\mathbf{x}) = \sum_{i \in \mathcal{V}} \psi_i(x_i) + \sum_{i \in \mathcal{V}, j \in \mathcal{N}_i} \psi_{ij}(x_i, x_j)$$

Data term **Smoothness term**

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \mathbf{L}} E(\mathbf{x})$$



Min Cut Problem – Image Segmentation

- Labeling problem : label pixels as foreground or background

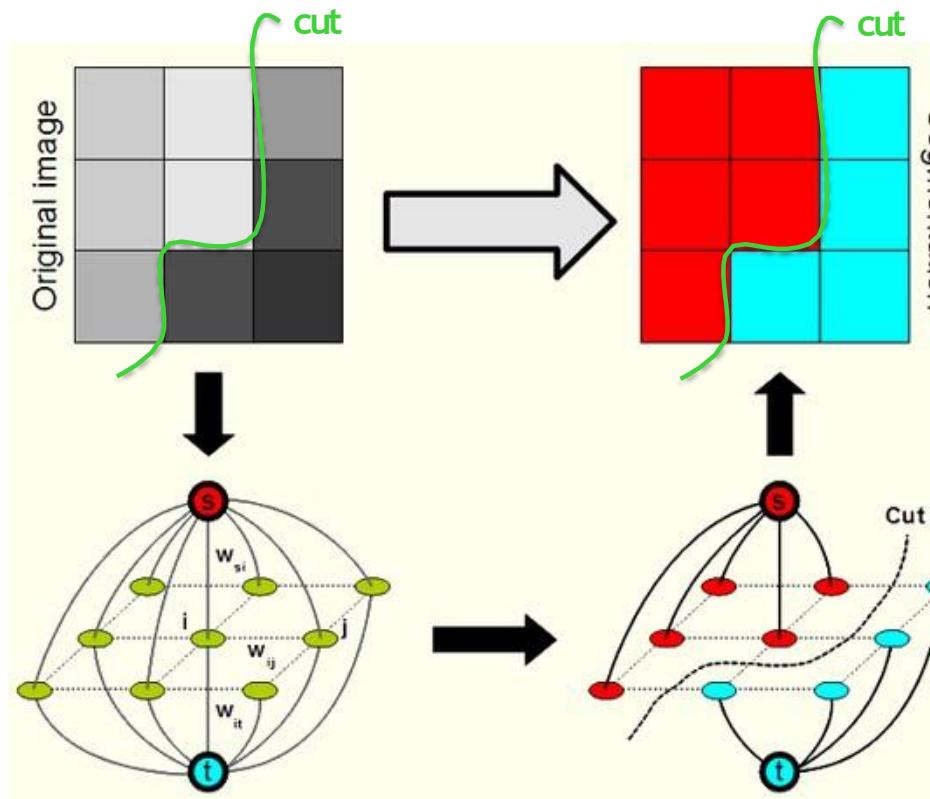


Image Source :<http://www.ondrej-danek.net/en/research>

For More Dedicated Study on Graph Cuts...

CVPR 2015 Tutorial on Energy Minimization and Discrete Optimization

Min Cut / Max Flow Algorithms, Alpha Expansion (multi label), etc.

https://inf.ethz.ch/personal/ladickyl/CVPR_Tutorial2015.htm

https://inf.ethz.ch/personal/ladickyl/gcut_cvpr15_part1.pdf

https://my.eng.utah.edu/~cs6320/cv_files/GraphCuts.pdf

Lubor Ladický

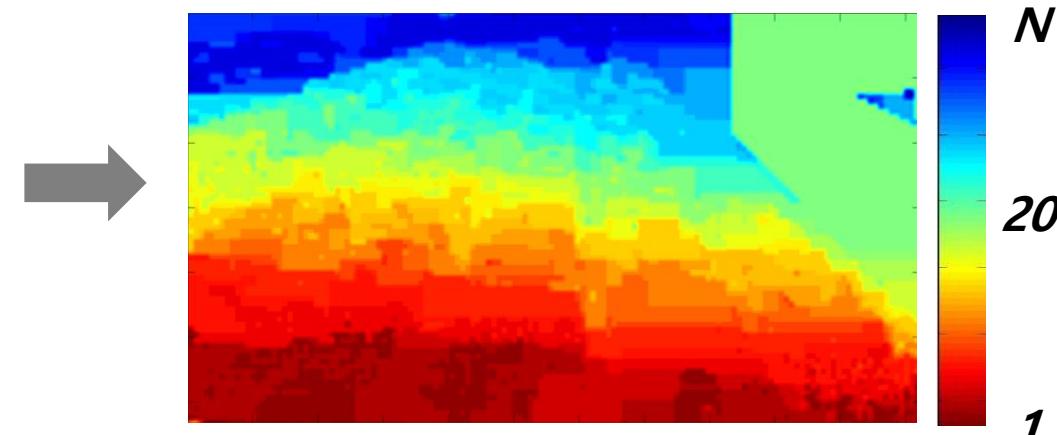


Depth(Index) Map Propagation Problem



Sparse depth map – confident (valid) pixels

[Multi-label Segmentation Problem]



Graph-Cut for Depth(Index) Map Propagation

$$E(\mathbf{x}) = \sum_{i \in S} \psi_i(x_i) + \sum_{i \in \mathcal{V}, j \in \mathcal{N}_i} \psi_{ij}(x_i, x_j)$$

Data term **Smoothness term**

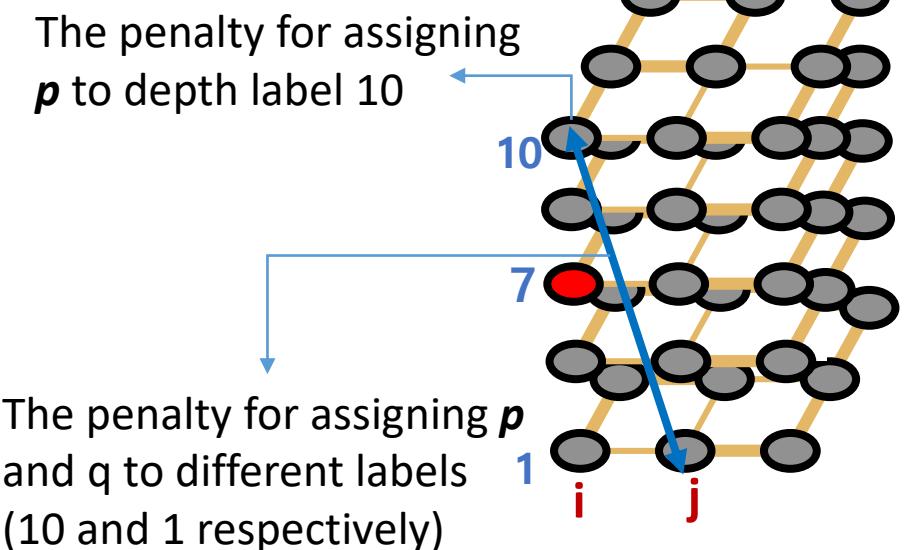
- **Data term** : $\psi_i(x_i)$

Key Idea. Assign close-by-value depth
for known data \mathbf{g}_i (from the initial depth)

- **Smoothness term** : $\psi_{ij}(x_i, x_j) = |x_i - x_j|$

Key Idea. Assign similar depth for neighbor pixels

Multi-label segmentation

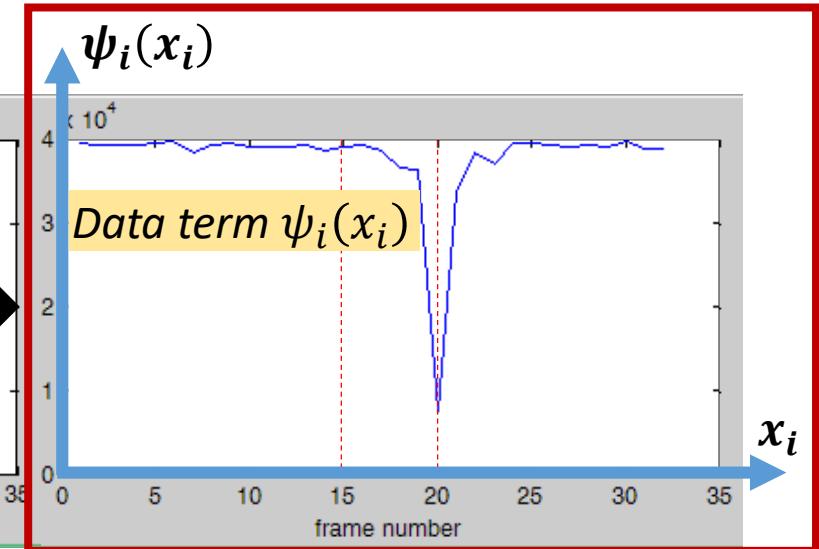
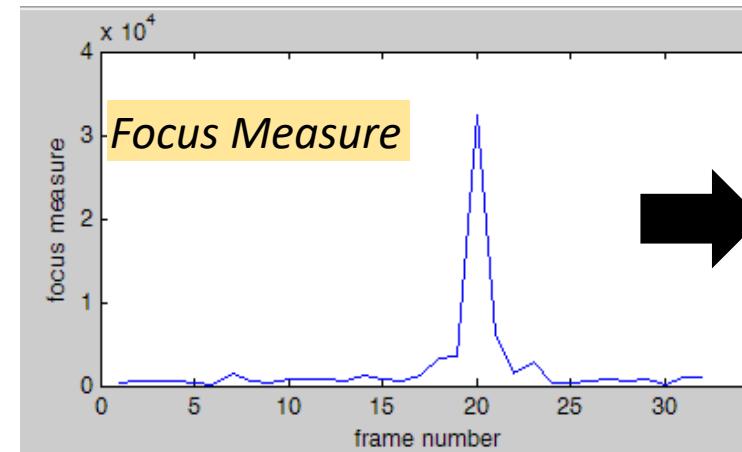


Another Energy Term Design

- **Data term** : $\psi_i(x_i)$ **Key Idea.** Assign close-by-value depth for known pixels
for each pixel:



Sparse Depth(Index) Map



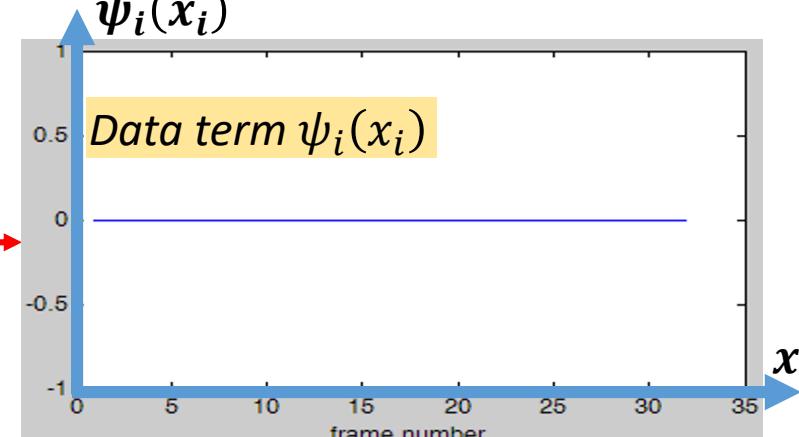
→ s.t. Pixel i prefers label 20

- Data term matrix `mat_D`

`mat_D` : $N \times H \times W$

N : # frames

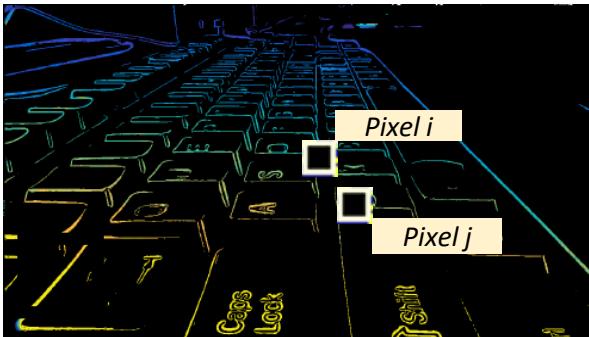
H, W : image height, width



→ s.t. Pixel i shows **flat** preferences

Another Energy Term Design

- **Smoothness term** : $\psi_{ij}(x_i, x_j)$ **Key Idea.** Assign similar depth for neighbor pixels
for each pixel pair (i, j) :



Sparse Depth(Index) Map

$$\psi_{ij}(x_i, x_j) = \begin{cases} |x_i - x_j| & \text{if } j \text{ is neighbor of } i \\ 0 & \text{Otherwise, ignore} \end{cases}$$

- Smoothness term matrix `mat_S`

mat S : N x N

N : # frames

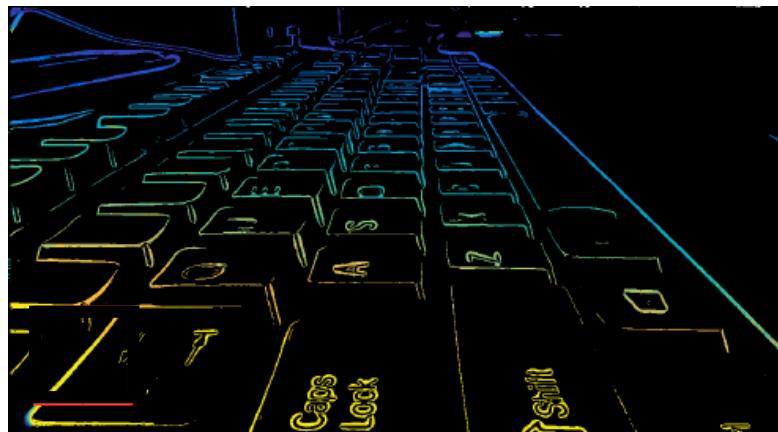
mat_S															
..							←	N	→					..	
↑	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
N	1	0	1	2	3	4	5	6	7	8	9	10	11	12	13
↓	2	1	0	1	2	3	4	5	6	7	8	9	10	11	12
..	3	2	1	0	1	2	3	4	5	6	7	8	9	10	11
..	4	3	2	1	0	1	2	3	4	5	6	7	8	9	10
..	5	4	3	2	1	0	1	2	3	4	5	6	7	8	9
..	6	5	4	3	2	1	0	1	2	3	4	5	6	7	8
..	7	6	5	4	3	2	1	0	1	2	3	4	5	6	7
..	8	7	6	5	4	3	2	1	0	1	2	3	4	5	6
..	9	8	7	6	5	4	3	2	1	0	1	2	3	4	5
..	10	9	8	7	6	5	4	3	2	1	0	1	2	3	4
..	11	10	9	8	7	6	5	4	3	2	1	0	1	2	3
..	12	11	10	9	8	7	6	5	4	3	2	1	0	1	2
..	13	12	11	10	9	8	7	6	5	4	3	2	1	0	1
..	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0

Graph-Cut for Depth(Index) Map Propagation

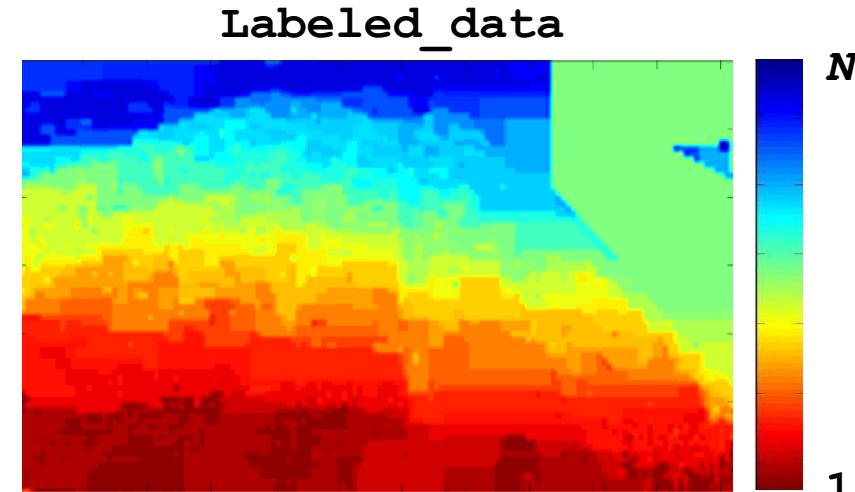
- **Optimization**

Using GCO-v3 MATLAB library

```
h = GCO_Create(H*W,N);  
GCO_SetDataCost(h,int32(mat_D));  
GCO_SetSmoothCost(h,int32(mat_S));  
:  
Labeled_data = GCO_GetLabeling(h); % (H*W)
```

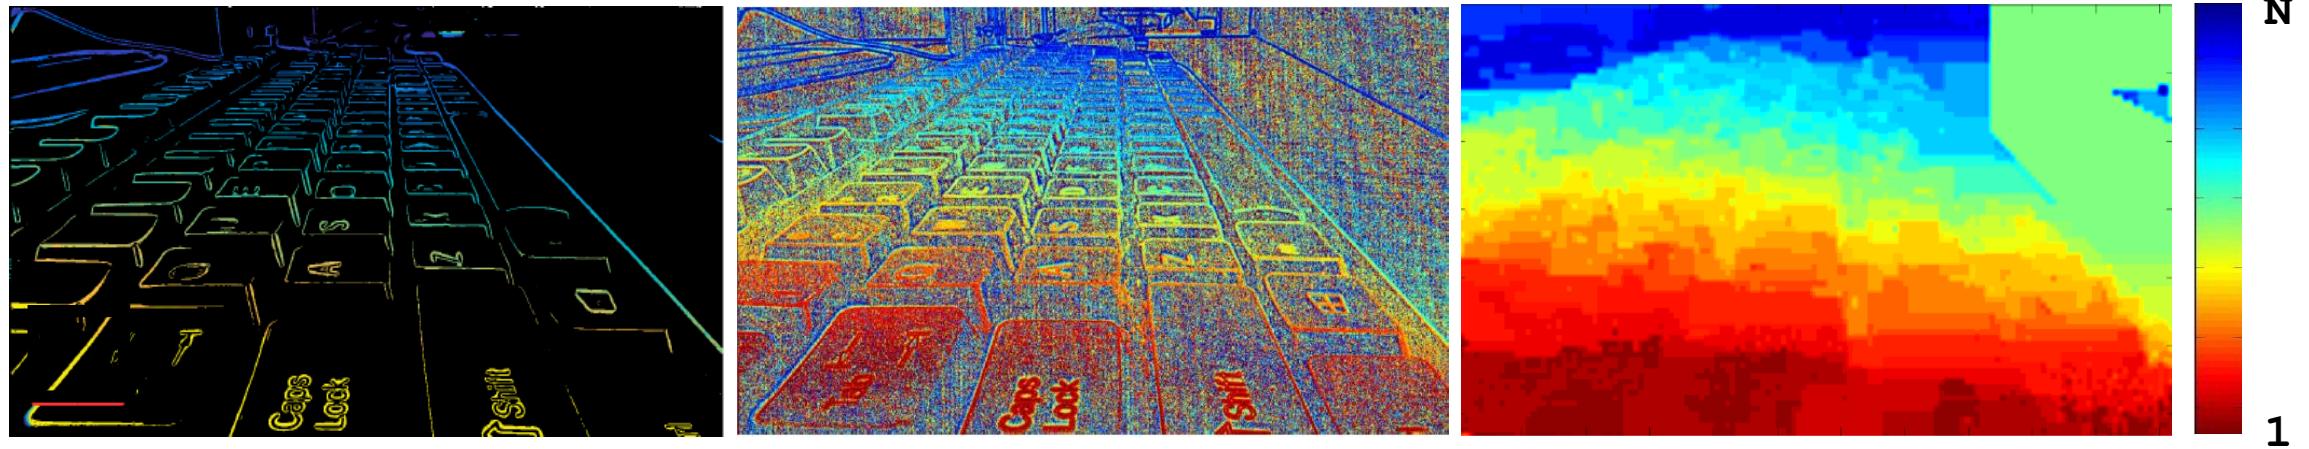


Sparse Depth (Index) map

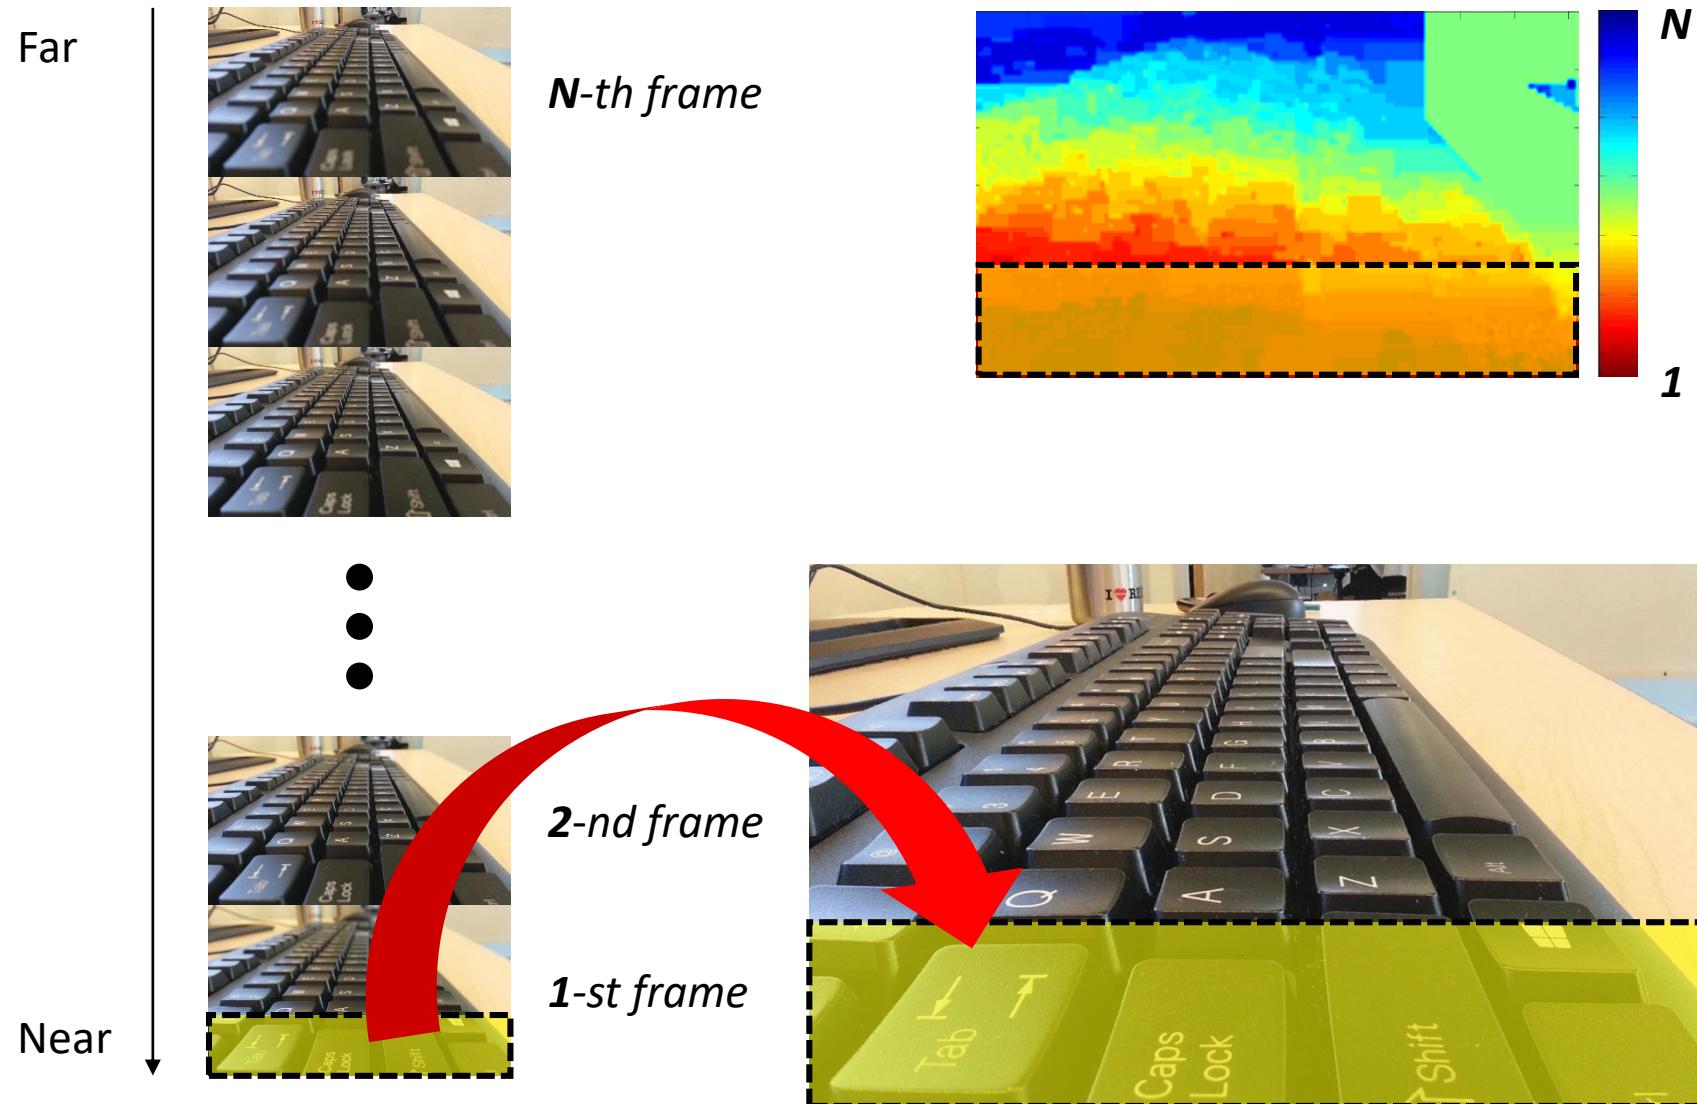


Graph-Cut Results

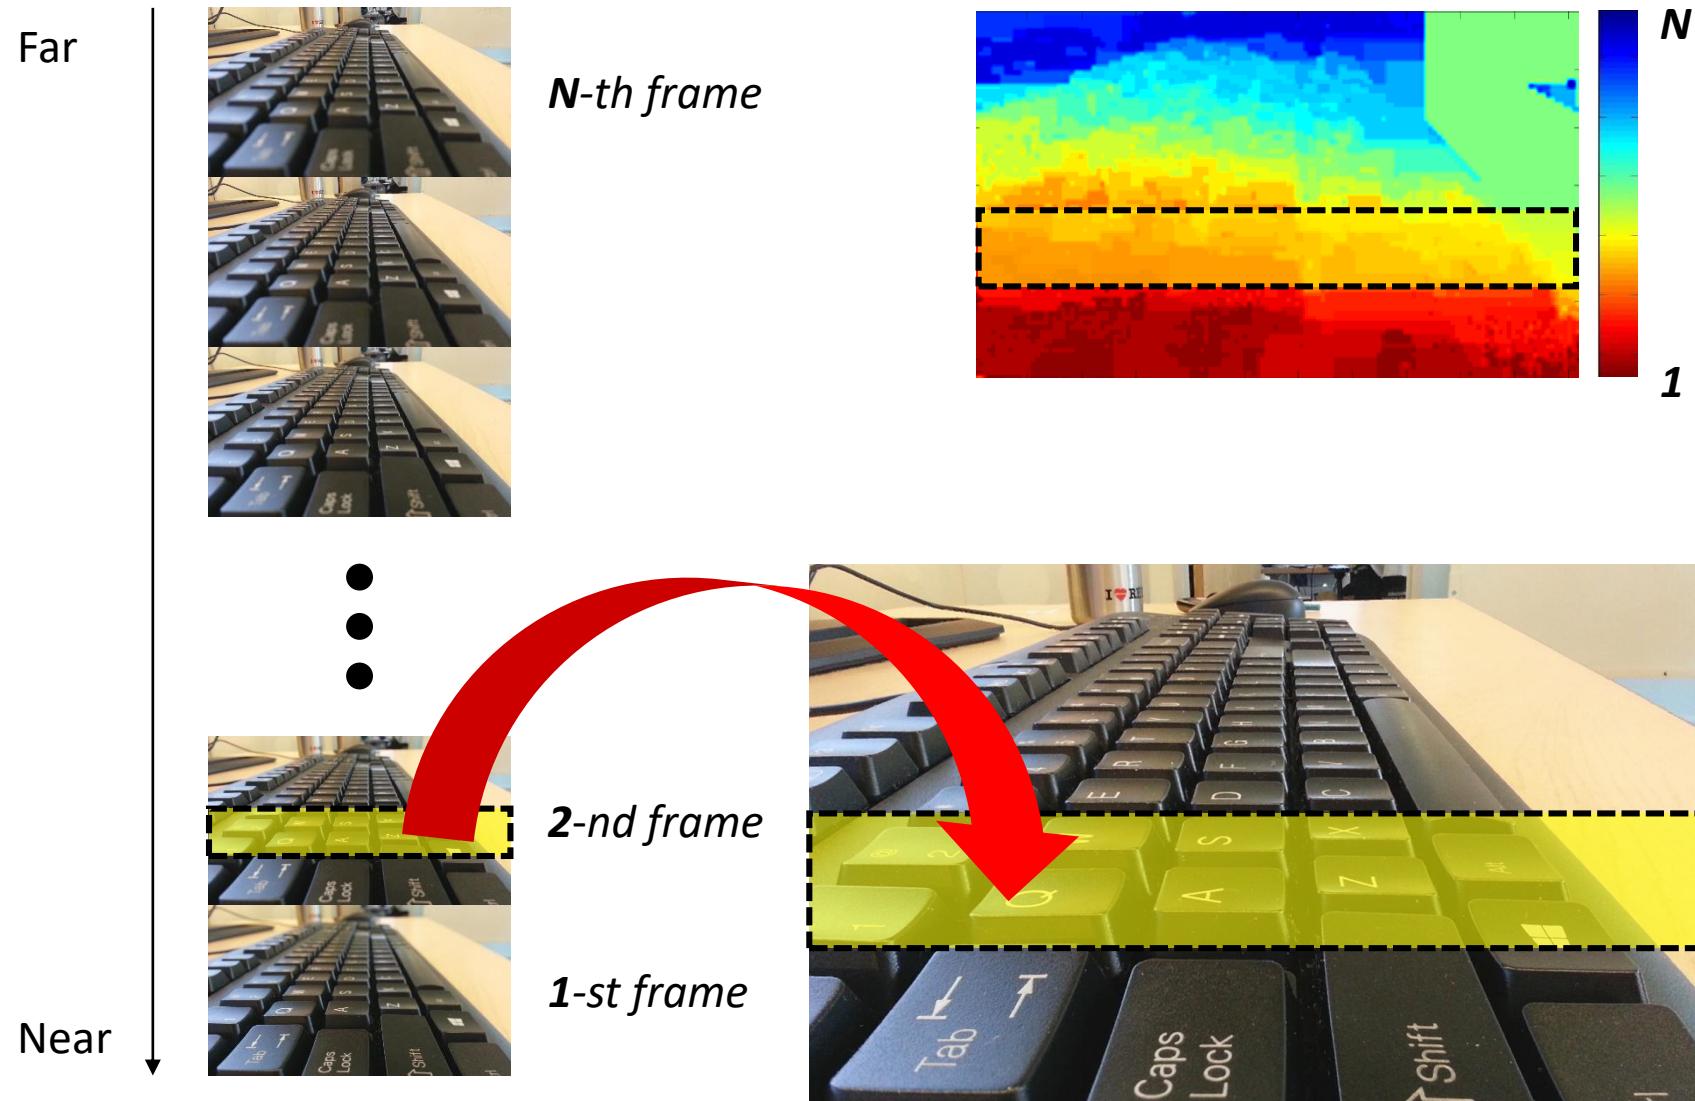
Sparse-to-Dense Depth Map Propagation



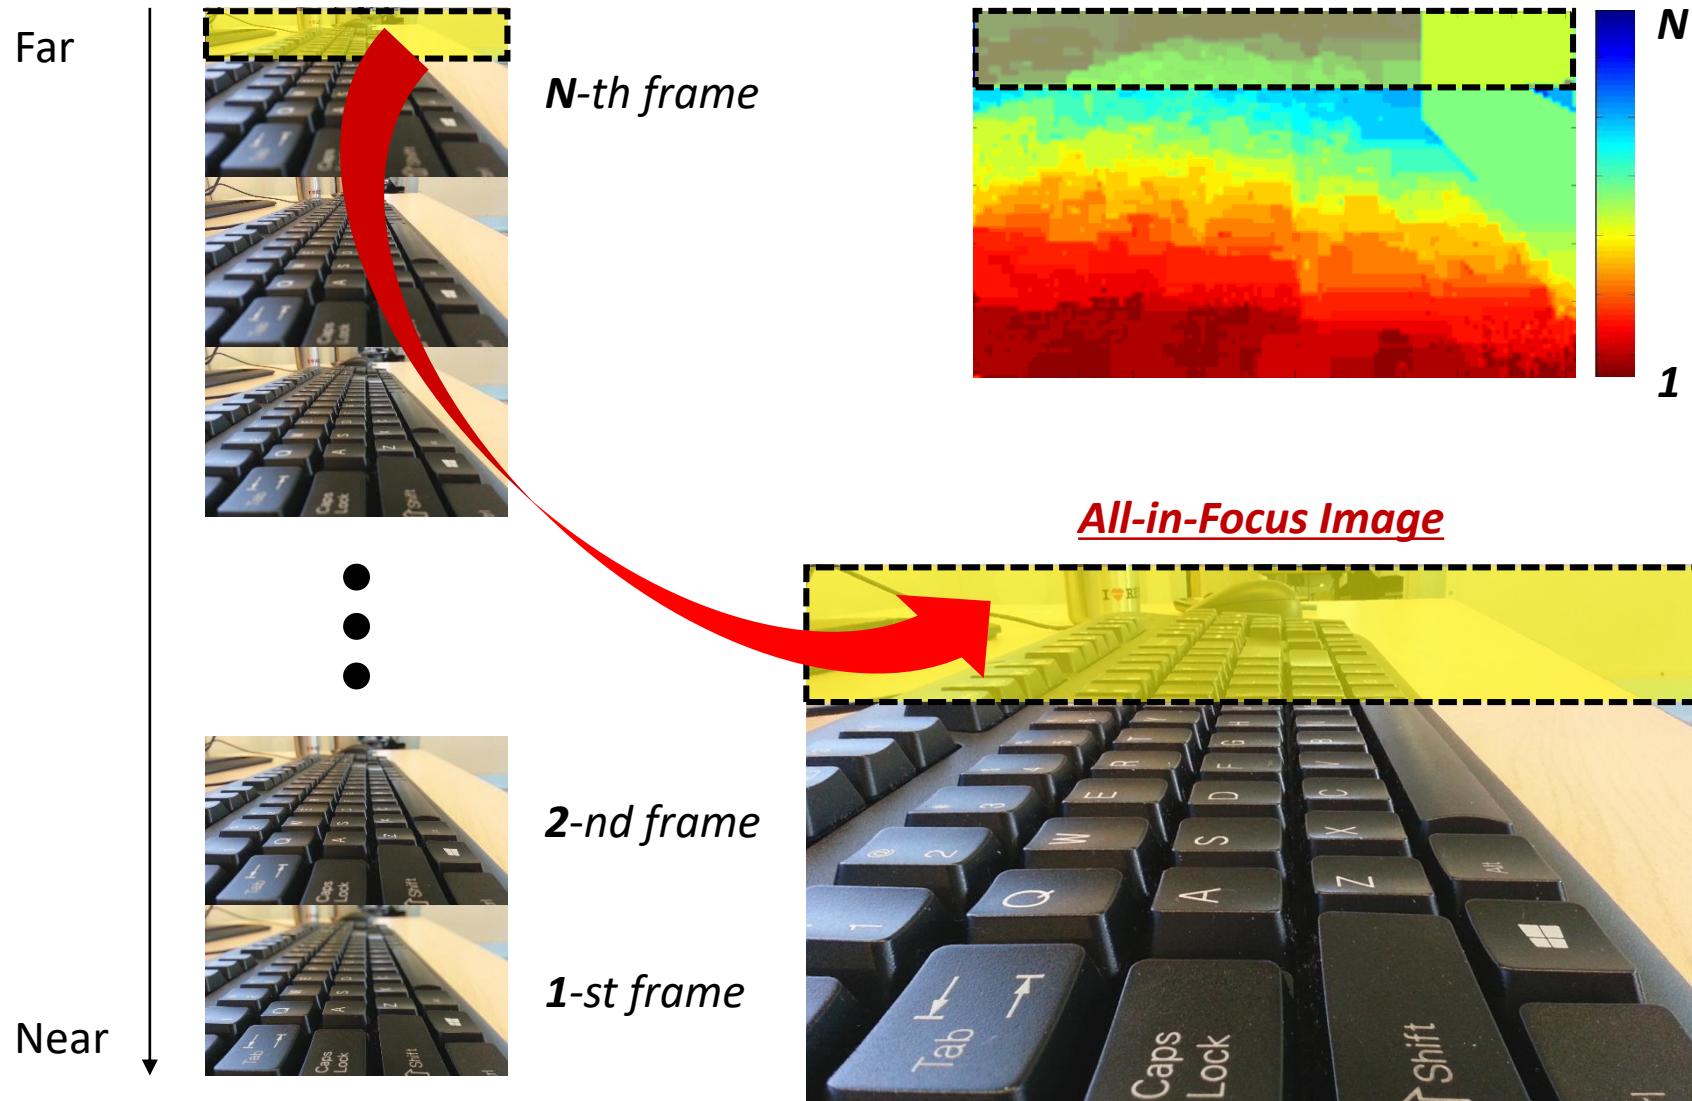
All-in-Focus Image Stitching



All-in-Focus Image Stitching

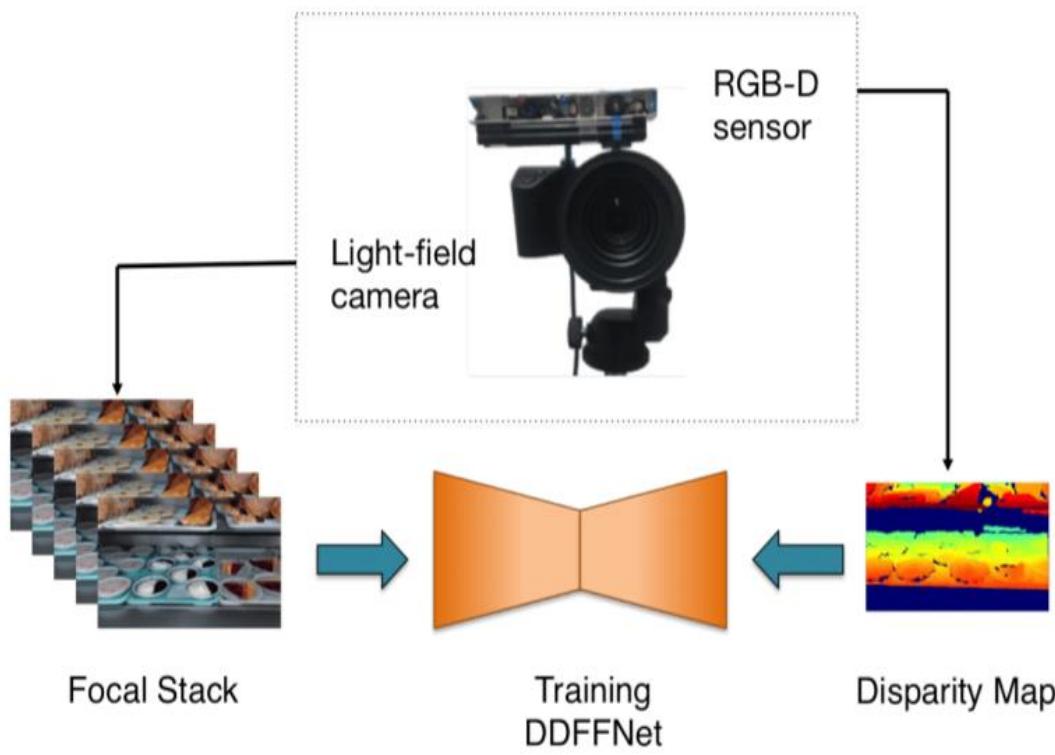


All-in-Focus Image Stitching



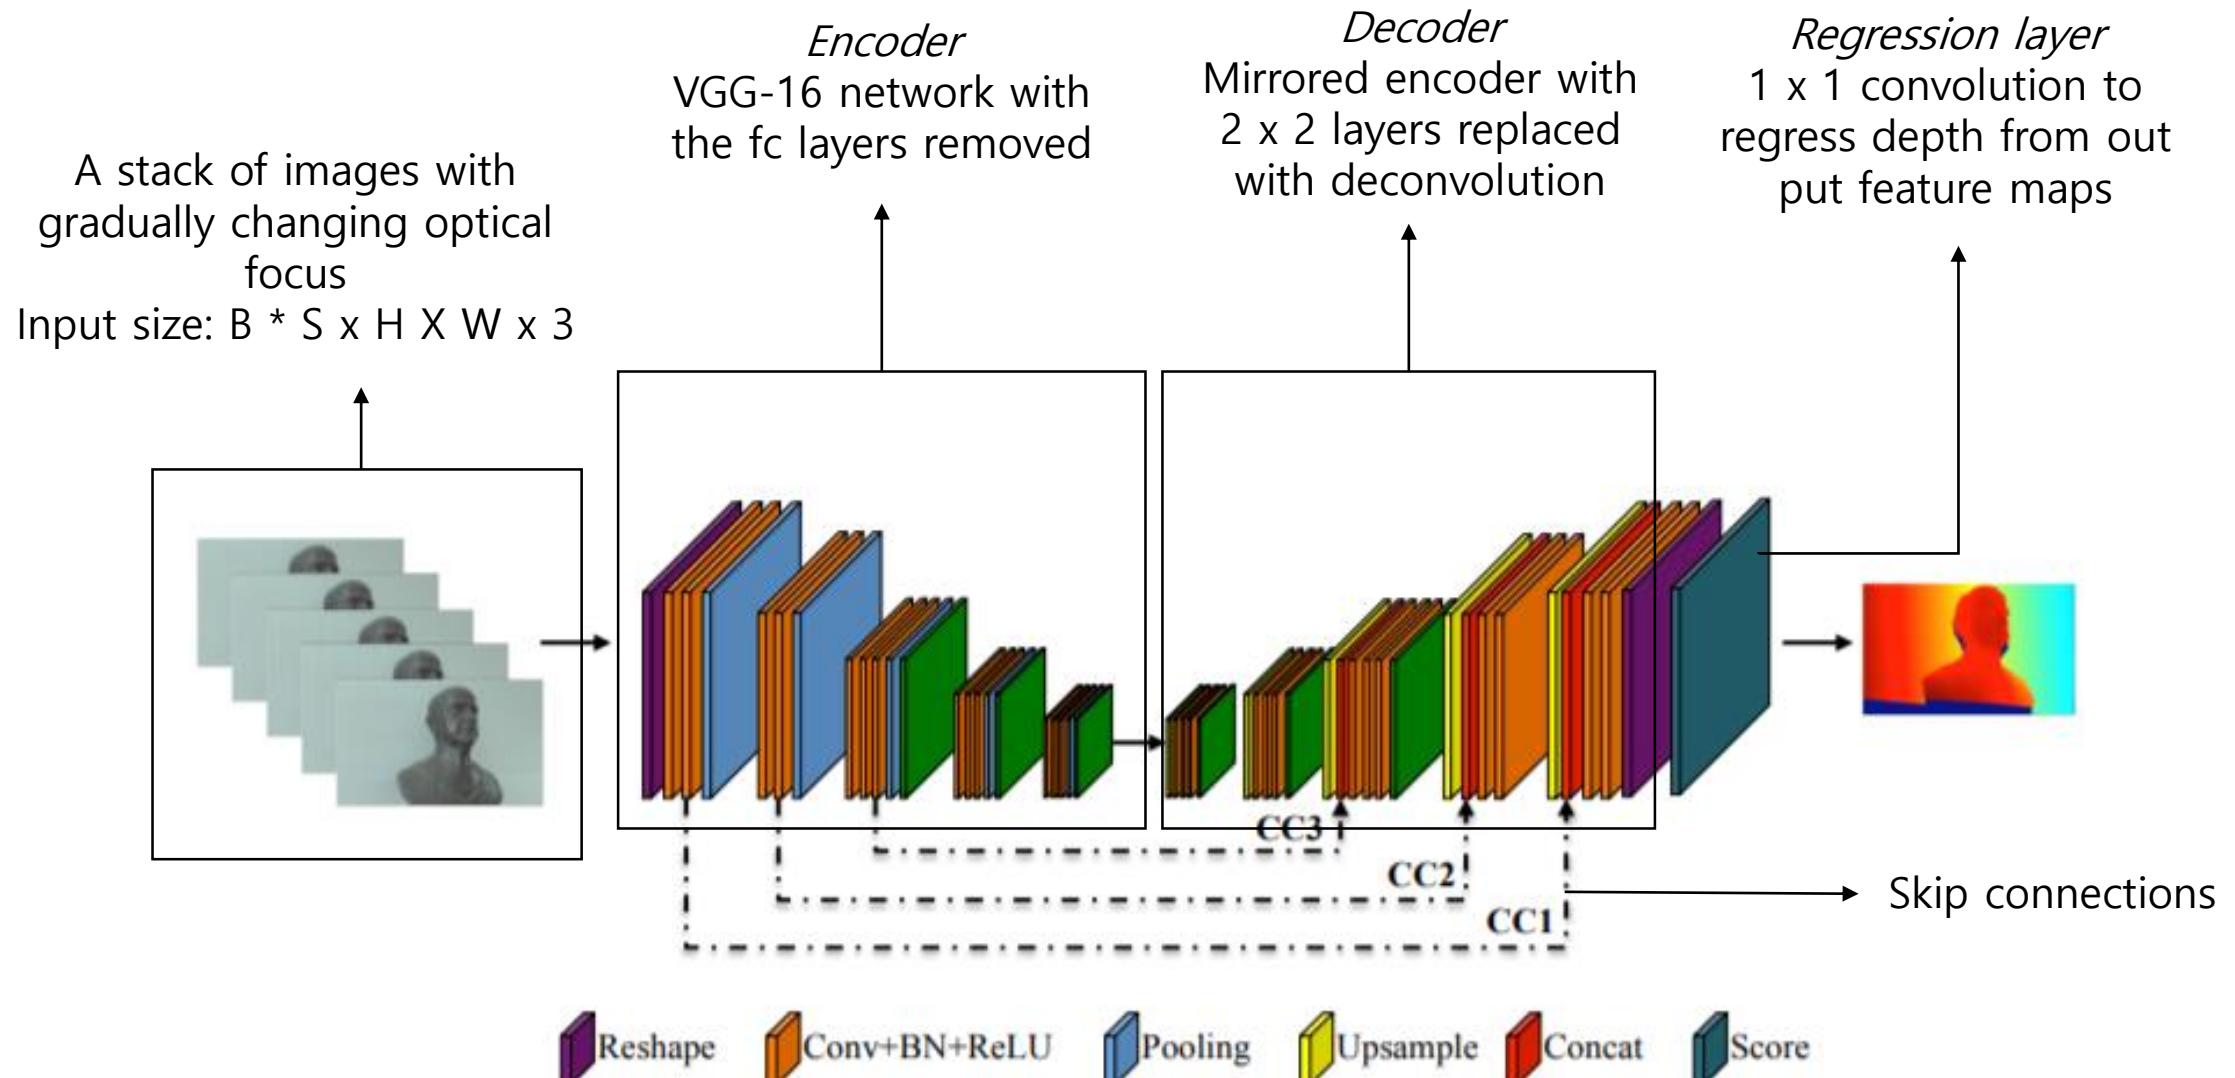
Network Architecture

- Use deep learning to reconstruct a pixel-accurate depth map from a given a stack of images with gradually changing optical focus



- Training DNNs require a significant amount of data.
- Manually obtaining a stack of refocused images is time consuming
- Light field imaging
 - ✓ One image per scene needs to be taken, meaning all images will have the same photographic exposure and the capturing process will be efficient.
 - ✓ Refocusing can be performed digitally, which allows to easily generate stacks with different focal steps
- RGB-D sensor for ground truth depth maps
- DDF 12-Scene benchmark
 - ✓ Dataset in 12 different indoor environments

Network Architecture



Quantitative Results

Method	MSE ↓	RMS ↓	log RMS ↓	Abs. rel. ↓	Sqr. rel. ↓	Accuracy ($\delta = 1.25$)				
						$\delta \uparrow$	$\delta^2 \uparrow$	$\delta^3 \uparrow$	Bump. ↓	
DDFFNet	Unpool	$2.9 e^{-3}$	0.050	0.50	0.64	0.05	39.95	63.46	78.26	0.62
	BL	$2.1 e^{-3}$	0.041	0.43	0.46	0.03	51.29	74.81	85.28	0.54
	UpConv	$1.4 e^{-3}$	0.034	0.33	0.30	0.02	52.41	83.09	93.78	0.54
	CC1	$1.4 e^{-3}$	0.033	0.33	0.37	0.02	60.38	82.11	90.63	0.75
	CC2	$1.8 e^{-3}$	0.039	0.39	0.39	0.02	44.80	76.27	89.15	0.75
	CC3	$9.7 e^{-4}$	0.029	0.32	0.29	0.01	61.95	85.14	92.99	0.59
	PSPNet	$9.4 e^{-4}$	0.030	0.29	0.27	0.01	62.66	85.90	94.42	0.55
	Lytro	$2.1 e^{-3}$	0.040	0.31	0.26	0.01	55.65	82.00	93.09	1.02
	PSP-LF	$2.7 e^{-3}$	0.046	0.45	0.46	0.03	39.70	65.56	82.46	0.54
	DFLF	$4.8 e^{-3}$	0.063	0.59	0.72	0.07	28.64	53.55	71.61	0.65
	VDFF	$7.3 e^{-3}$	0.080	1.39	0.62	0.05	8.42	19.95	32.68	0.79

Table 1: **Quantitative results.** DDFFNet-CC3 is the best depth from focus method and provides also better results compared to Lytro, *i.e.* depth from light-field. Metrics are computed on the predicted and the groundtruth disparity maps.

- ❖ DDFFNet is more accurate by a large margin when compared to classical variational DFF method.
- ❖ Faster on a GPU (orders of magnitude)

VDFF: Variational Depth from Focus Reconstruction, TIP, 2015

PSPNet: Pyramid Scene Parsing Network, CVPR, 2017

- MSE : $\frac{1}{|\mathcal{M}|} \sum_{p \in \mathcal{M}} \|f(\mathcal{S}_p) - D_p\|_2^2$
- RMS : $\sqrt{\frac{1}{|\mathcal{M}|} \sum_{p \in \mathcal{M}} \|f(\mathcal{S}_p) - D_p\|_2^2}$
- log RMS : $\sqrt{\frac{1}{|\mathcal{M}|} \sum_{p \in \mathcal{M}} \|\log f(\mathcal{S}_p) - \log D_p\|_2^2}$
- Absolute relative : $\frac{1}{|\mathcal{M}|} \sum_{p \in \mathcal{M}} \frac{|f(\mathcal{S}_p) - D_p|}{D_p}$
- Squared relative : $\frac{1}{|\mathcal{M}|} \sum_{p \in \mathcal{M}} \frac{\|f(\mathcal{S}_p) - D_p\|_2^2}{D_p}$
- Accuracy: % of D_p s.t $\max\left(\frac{f(\mathcal{S}_p)}{D_p}, \frac{D_p}{f(\mathcal{S}_p)}\right) = \delta < thr$
- BadPix(τ): $\frac{|\{p \in \mathcal{M} : |f(\mathcal{S}_p) - D_p| > \tau\}|}{|\mathcal{M}|} \cdot 100$
- Bumpiness: $\frac{1}{|\mathcal{M}|} \sum_{p \in \mathcal{M}} \min(0.05, \|\mathbf{H}_\Delta(p)\|_F) \cdot 100$

Deep Depth from Defocus

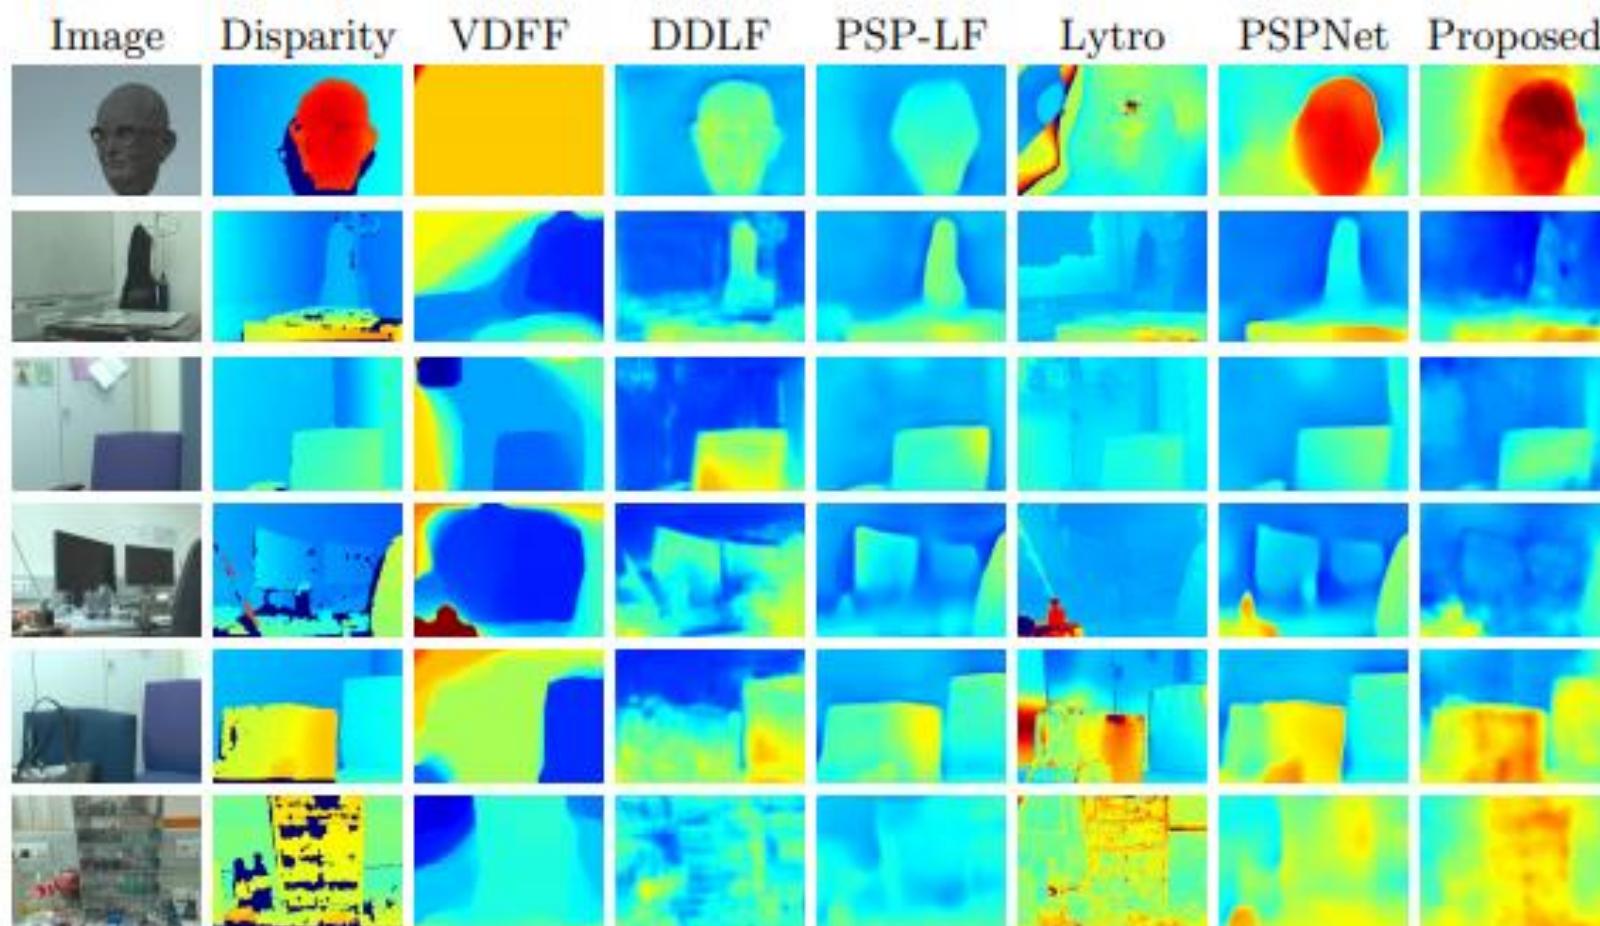


Fig. 6: Qualitative results of the DDFFNet versus state-of-the-art methods. Results are normalized by the maximum disparity. Warmer colors represent closer distances. Best viewed in color.

Deep Depth from Defocus

- What is the network learning?
 - Giving the stack of image to the network at the same time helps the network identify which pixel is the sharpest
 - The regression layer then regresses the depth from sharpness (focus)

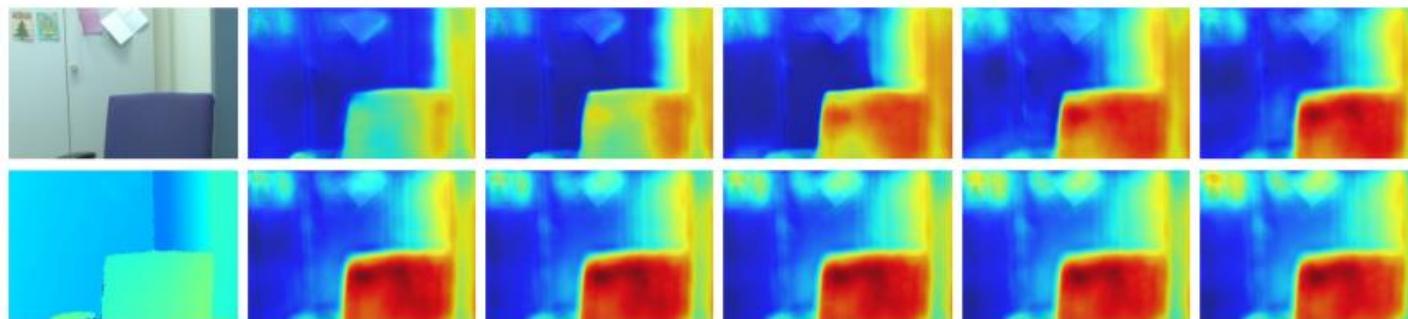


Fig. 2: **Activation heat maps** for the refocused images in a focal stack. First column: *top*: center sub-aperture image, *bottom*: groundtruth disparity map. The rest of the columns show from left to right, top to bottom, how the activations on the focal stack images evolve. Note how the activations slowly shift to the closest object (the chair) as we advance in the focal stack.

- Code and dataset are publicly available: <https://github.com/soyers/ddff-pytorch>
- More details can be found here: <https://hazirbas.com/projects/ddff/>